




Popularity of Algorithms



Today we'll cover

- Linear Regression
- Logistic Regression
- Regularized Regression
- Decision Tree
- Random Forests (Bagging)
- K-Nearest Neighbours
- Ensemble Model

<https://www.kdnuggets.com/2019/04/top-data-science-machine-learning-methods-2018-2019.html>



No Free Lunch

No Free Lunch แปลว่า “ไม่มีโมเดลไหนเก่งที่สุด และสามารถตอบโจทย์ได้ทุกปัญหา”

ถ้ามีใครถามว่าโมเดลไหนเก่งที่สุด?

ให้ตอบว่า “It depends” (ขึ้นอยู่กับข้อมูล)

ความท้าทายของ ML คือการหาโมเดลที่ดีที่สุดสำหรับปัญหาที่เรากำลังแก้

R Occam's Razor

Algorithm #1

vs.

Algorithm #2

ถ้ามีโมเดลสองตัวที่มี performance ดีเท่าๆกัน ให้เลือกตัวที่สร้างและอธิบายได้ง่ายกว่า (**choose simpler model**)



How to choose a model


ให้ลองถาม 2 คำถามง่ายๆนี้

1. ปัญหานี้เป็น regression หรือ classification?
2. อยากได้ high accuracy หรือ high interpretability?
 - Always choose a simpler model if performances are similar
 - Try different algorithms and find the right one.



Caret Package

Learn more at <https://topepo.github.io/caret/index.html>



Max Kuhn

the author of caret package





Dataset for our projects

```
## load library
## install.packages("mlbench")
library(mlbench)
library(tidyverse)

## load dataset for regression
data("BostonHousing")
glimpse(BostonHousing)

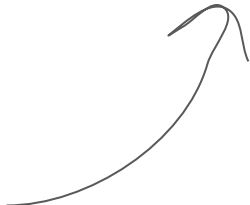
## load dataset for classification
data("PimaIndiansDiabetes")
glimpse(PimaIndiansDiabetes)
```





Caret Training Template

```
model <- train(form = y ~ . ,  
               data = train_data ,  
               method = "lm" )
```

Model that we
want to train




R Our first machine



R Euclidean Distance

$$d = \sqrt{(x1 - x2)^2 + (y1 - y2)^2}$$




Euclidean Distance

$$d = \sqrt{(x1 - x2)^2 + (y1 - y2)^2}$$


```
point_1 <- c(2,3)
point_2 <- c(6,8)
d <- sqrt( (2-6)**2 + (3-8)**2 )
print(d)
```



We use majority vote to assign label



Predict 'Red' $3/5 = 60\%$



Predict 'Green' $3/5 = 60\%$

Majority Vote




Steps to train this algorithm

1. แต่งตั้ง สว 250 คน
2. ช่วยกันเลือกนายกฯ

เฮ้ย เต๋วๆๆๆ 555555555555+

R We use average value for regression problem



Use the average as prediction

$$(10 + 16 + 18 + 20 + 15) / 5 = 15.8$$



K-Nearest Neighbors

1. Choose K
2. Compute distance
3. Majority vote for classification or
Average for regression

Train test split (the easiest method)

Prepare dataset first

We'll use split data into training 75% and testing 25%

```
## split data
set.seed(99)
n <- nrow(BostonHousing)
id <- sample(n, size = n*0.75, replace=FALSE)
train_data <- BostonHousing[id, ]
test_data <- BostonHousing[-id, ]
```





Very easy to train a machine in R

```
## train model
set.seed(99)
knn_model <- train(medv ~ .,
                   data = train_data,
                   method = "knn")

## test model
p <- predict(knn_model, newdata = test_data)

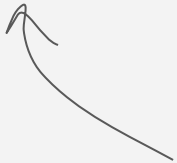
## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```




```
## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
                  data = train_data,
                  method = "knn",
                  trControl = ctrl)

## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```



5 Fold Cross Validation

Random Search

```
## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
  data = train_data,
  tuneLength = 5,
  method = "knn",
  trControl = ctrl)
```

Try 5 values of K



```
## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```



Grid Search

```
## create grid
myGrid <- expand.grid(k = 1:10)

## train model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5, verboseIter = TRUE)
knn_model <- train(medv ~ .,
  data = train_data,
  tuneGrid = myGrid,
  method = "knn",
  trControl = ctrl)

## test model
p <- predict(knn_model, newdata = test_data)

## rmse
rmse <- sqrt(mean((p - test_data$medv)**2))
```

The values we
select ourselves :D



R

Grid Search Result

```
> knn_model
k-Nearest Neighbors

379 samples
13 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 303, 303, 303, 303, 304
Resampling results across tuning parameters:
```

k	RMSE	Rsquared	MAE
1	7.607647	0.4598462	4.922035
2	6.857741	0.5208036	4.750591
3	6.778822	0.5139722	4.657967
4	6.659557	0.5153593	4.651180
5	6.646851	0.5131619	4.681624
6	6.660705	0.5081547	4.648119
7	6.711653	0.5001402	4.661983
8	6.881981	0.4749499	4.749852
9	6.872293	0.4768345	4.763948
10	6.927021	0.4683690	4.773996

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was k = 5.

Cross Validation
ช่วยเราเลือกค่า k ที่
ทำให้ RMSE ต่ำที่สุด
ตอนเรา train model

1. KNN เข้าใจง่ายทำงานได้โอเคร ถ้า feature ไม่เยอะมาก
2. KNN ใช้ได้ทั้ง regression/ classification
3. K ใน KNN คือค่า hyperparameter ที่เราเปลี่ยนได้
4. เราเลือก K ที่ทำให้ train RMSE ต่ำที่สุด
5. train RMSE ต่ำที่สุดไม่ได้แปลว่าโมเดลเราจะทำนาย test_data ได้ดี ต้องเอาไปทดสอบอีกที



Caret Interface Summary



Classification vs. Regression

Classification

```
set.seed(42)

ctrl <- trainControl(method = "cv",
                     number = 5)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "Accuracy",
  trControl = ctrl
)
```


Regression

```
set.seed(42)

ctrl <- trainControl(method = "cv",
                     number = 5)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "RMSE",
  trControl = ctrl
)
```

Same interface, different metrics





Classification Interfaces

Classification - ROC Sens Specs

```
set.seed(42)

ctrl <- trainControl(
  method = "cv",
  number = 5,
  summaryFunction = twoClassSummary,
  classProbs = TRUE)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "ROC",
  trControl = ctrl
)
```

Classification - AUC Precision Recall F1


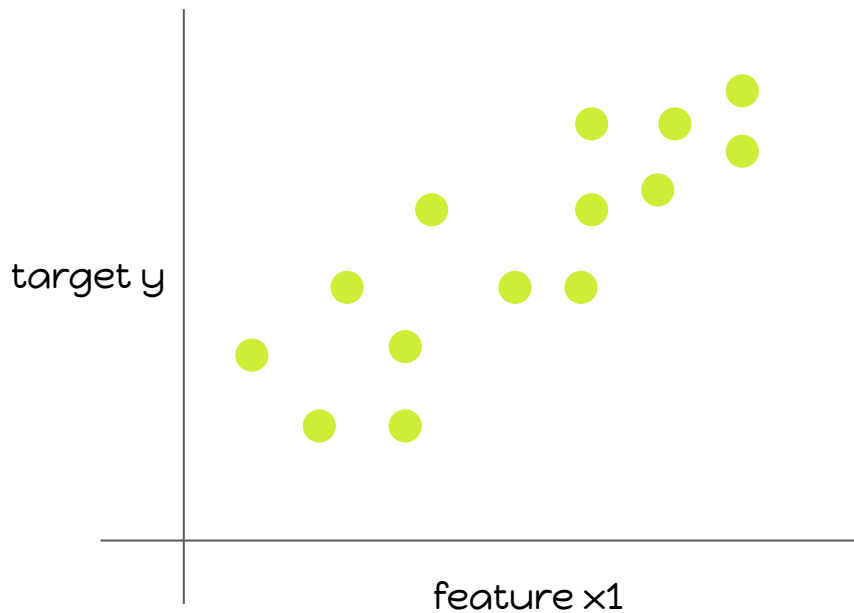
```
set.seed(42)

ctrl <- trainControl(
  method = "cv",
  number = 5,
  summaryFunction = prSummary,
  classProbs = TRUE)

model <- train(
  y ~ .,
  data = df,
  method = "knn",
  metric = "AUC",
  trControl = ctrl
)
```




Linear Regression Explained




Linear Regression Model

$$\hat{y} = B_0 + B_1 X_1$$

prediction

y intercept

slope



Linear Regression finds
B0 and B1 that minimize
the error

Minimize Error

$$\text{minimize } \sum (\text{prediction} - \text{actual})^2$$

$$\text{minimize } \sum (\hat{y} - y)^2$$



Sum of Squared Error
or RSS (for short)

R

Common Regression Metrics

$$MAE = \frac{1}{n} * \sum |\hat{y} - y|$$

$$MSE = \frac{1}{n} * \sum (\hat{y} - y)^2$$

$$RMSE = \sqrt{\frac{1}{n} * \sum (\hat{y} - y)^2}$$

โมเดลที่เราเทรนจะพยายามทำให้ค่า
MAE/ MSE/ RMSE มีค่าต่ำที่สุด
i.e. minimize error



Easy to compute in Spreadsheets

y	y_hat	error	error	error^2	
10	8.5	1.5	1.5	2.25	
12	14.5	-2.5	2.5	6.25	
14	10	4	4	16	
16	17	-1	1	1	
18	17.5	0.5	0.5	0.25	
			9.5	25.75	
			1.9	5.2	2.3
			MAE	MSE	RMSE



Build linear regression in R

```
## train model with train_data
set.seed(99)
lm_model <- train(medv ~ rm + indus + crim,
                  data = train_data,
                  method = "lm")

## test model (predict test data)
p <- predict(lm_model, newdata = test_data)
rmse <- sqrt(mean( (p - test_data$medv)** 2 ))
```





Linear regression with K-Fold

```
## train model with train_data
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                     verboseIter = TRUE)

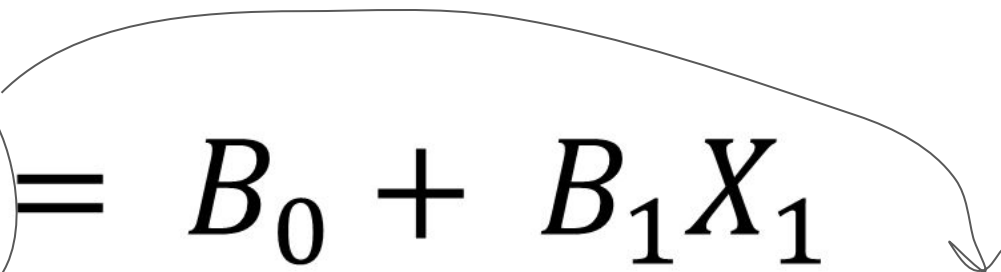
lm_model <- train(medv ~ rm + indus + crim,
                 data = train_data,
                 method = "lm",
                 trControl = ctrl)

## test model (predict test data)
p <- predict(lm_model, newdata = test_data)
rmse <- sqrt(mean( (p - test_data$medv)** 2 ))
```

R Logistic regression for binary classification



R Logistic is very similar to linear regression

$$Z = B_0 + B_1 X_1$$



$$P(Y = 1|x) = \frac{e^z}{1 + e^z}$$

Sigmoid function



Z	sigmoid(Z)	y_hat
-10	0.0000	0
-9	0.0001	0
-8	0.0003	0
-7	0.0009	0
-6	0.0025	0
-5	0.0067	0
-4	0.0180	0
-3	0.0474	0
-2	0.1192	0
-1	0.2689	0
0	0.5000	0
1	0.7311	1
2	0.8808	1
3	0.9526	1
4	0.9820	1
5	0.9933	1
6	0.9975	1
7	0.9991	1
8	0.9997	1
9	0.9999	1
10	1.0000	1



If $\text{sigmoid}(Z) > 0.5$, predict $y = 1$, else $y = 0$



R Our prediction changes if threshold changes



R Code

```
## train model with train_data
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                     verboseIter = TRUE)

logistic_model <- train(diabetes ~ .,
                       data = train_data,
                       method = "glm",
                       trControl = ctrl)

## test model (predict test data)
p <- predict(logistic_model, newdata = test_data)
accuracy <- mean(p == test_data$diabetes)
```



- Accuracy
- Precision
- Recall
- F1




สามารถคำนวณได้
ง่าย ๆ จาก

Confusion Matrix

R

Common classification metrics



AUC

Area Under Curve



Interpretation in Thai

Metrics	ความหมาย
Accuracy	ความถูกต้องของโมเดลในภาพรวม
Precision	ทุก 100 ครั้งที่เรากำหนด $y=1$ โอกาสถูกเท่าไร
Recall	ทุกผู้ป่วยจริงๆ 100 คน เราตรวจเจอที่คน
F1	ค่าเฉลี่ยระหว่าง precision, recall

27.4.5 Balanced accuracy and F_1 score

Although we usually recommend studying both specificity and sensitivity, very often it is useful to have a one-number summary, for example for optimization purposes. One metric that is preferred over overall accuracy is the average of specificity and sensitivity, referred to as *balanced accuracy*. Because specificity and sensitivity are rates, it is more appropriate to compute the *harmonic average*. In fact, the F_1 -score, a widely used one-number summary, is the harmonic average of precision and recall:

$$\frac{1}{\frac{1}{2} \left(\frac{1}{\text{recall}} + \frac{1}{\text{precision}} \right)}$$

Because it is easier to write, you often see this harmonic average rewritten as:

$$2 \times \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

when defining F_1 .



R Code - Confusion Matrix

```
## use table()
table(predicted, actual, dnn = c("predicted", "actual"))
```

		Actual	
		neg	pos
Predicted	neg	101	33
	pos	14	44

```
## how we calculate four metrics
accuracy <- (101 + 44) / (101 + 33 + 44 + 14)
precision <- 44 / (44 + 14)
recall <- 44 / (44 + 33)
F1 <- 2 * (precision * recall) / (precision + recall)
```



Regularized Regression

1. Ridge Regression
2. Lasso Regression

Regularization is a key technique in ML to **reduce overfitting** :D

R

Ridge Regression (L2)

$$RSS = \sum (\hat{y} - y)^2$$

Normal RSS from Linear Regression

$$Ridge\ RSS = \sum (\hat{y} - y)^2 + \lambda \sum \beta^2$$

Ridge add this term to the error function

R

Lasso Regression (L1)

$$RSS = \sum (\hat{y} - y)^2$$

Normal RSS from Linear Regression

$$Lasso\ RSS = \sum (\hat{y} - y)^2 + \lambda \sum |\beta|$$

Lasso add this term to the error function



Regularization helps reduce overfitting

$$y_hat = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4$$

$$y_hat = 100 + 150x_1 + 200x_2 + 120x_3 + 80x_4$$


$$\text{Lasso} = \text{RSS} + \text{Lambda} * (150 + 200 + 120 + 80)$$

$\text{Lambda} = 0$

Error is the same as Linear Regression

$\text{Lambda} > 0$

All coefficient (B) in the model must be **shrunk** to reduce the new error

script.R

```
1 # Train glmnet with custom trainControl and tuning: model
2 model <- train(
3   y ~ .,
4   data = overfit,
5   tuneGrid = expand.grid(
6     alpha = 0:1,
7     lambda = seq(0.0001, 1, length=20)
8   ),
9   method = "glmnet",
10  trControl = myControl
11 )
12
13 # Print model to console
14 model
15
16 # Print maximum ROC statistic
17 max(model$results$ROC)
```

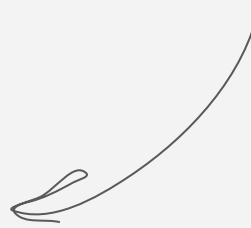


ElasticNet model

```
## train elasticnet model
set.seed(99)
ctrl <- trainControl(method = "cv", number = 5,
                     verboseIter = TRUE)

enet_model <- train(diabetes ~ .,
                   data = train_data,
                   method = "glmnet",
                   trControl = ctrl)
```

ElasticNet =
Mixed between
Ridge + Lasso



```
## test model
p <- predict(enet_model, newdata = test_data)
accuracy <- mean(p == test_data$diabetes)
```





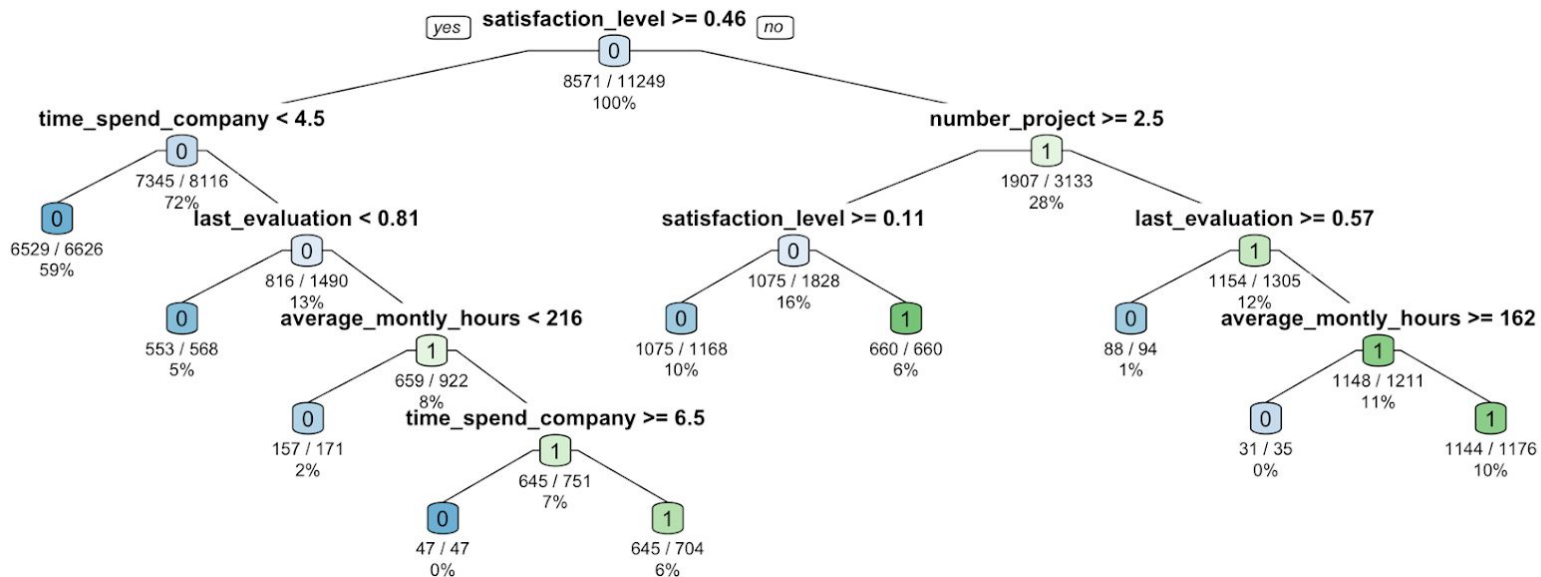
Time for a fun game :D

- Ask me a yes/no question
- To guess my favourite animal
- Max 10 questions











Decision Tree





How decision tree work?







Age	Sex	App
15	M	 tinder
20	M	 tinder
16	F	
19	M	 tinder
22	F	
20	F	

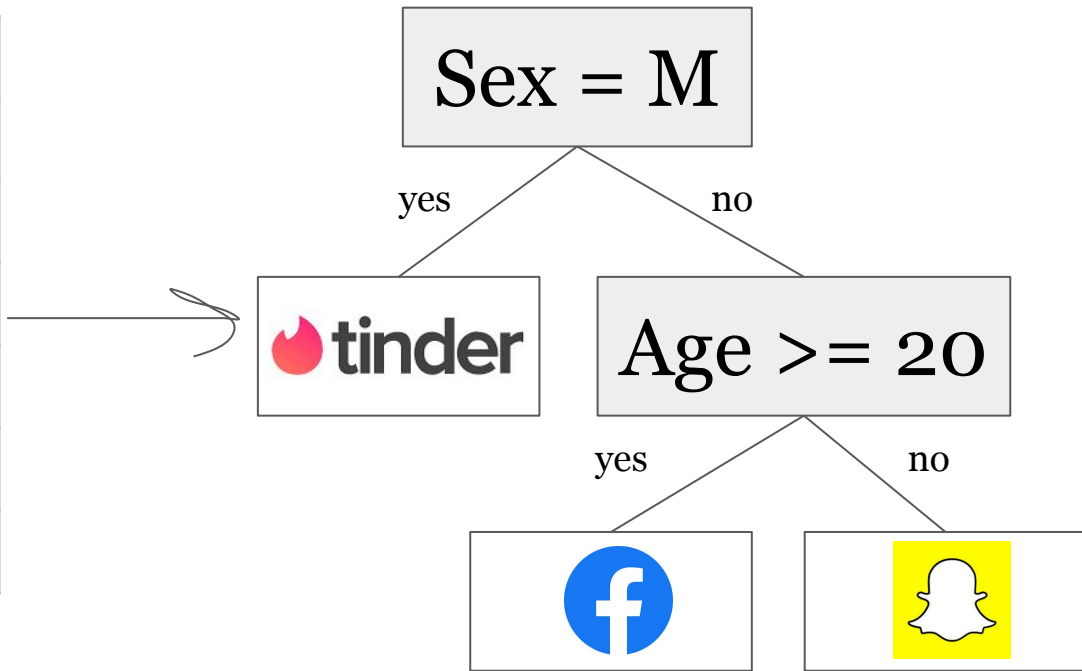
เวลาเราสร้าง decision tree เราถามคำถาม yes/no question ทีละข้อ

i.e. feature ใช้แบ่ง App ได้ดีที่สุด



How decision tree work?

Age	Sex	App
15	M	 tinder
20	M	 tinder
16	F	
19	M	 tinder
22	F	
20	F	



Decision Tree with K-Fold

```
## train tree
set.seed(99)

ctrl <- trainControl(method = "...", number = ...,
                     verboseIter = TRUE)

tree_model <- train(diabetes ~ .,
                   data = ...,
                   method = "rpart",
                   trControl = ...)

## test model
p <- predict(tree_model, newdata = ...)
accuracy <- mean(...)
```



Let's do a quick review

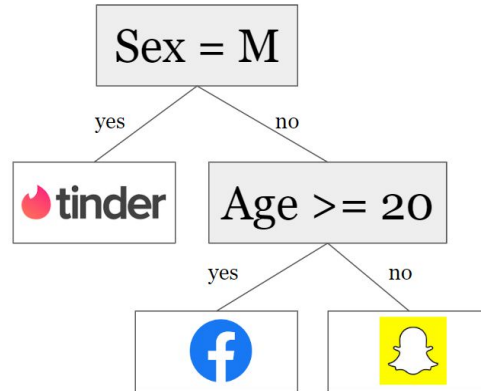

Parametric	Non-Parametric
Linear Regression	KNN
Logistic Regression	Decision Tree
Ridge Regression	Random Forest
Lasso Regression	

Regression is a Linear Combination (Parametric)

$$\hat{y} = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4$$

↖
มีพจน์

While a non-parametric has no form :P



Random Forest



We grow hundreds of
uncorrelated (decision) trees





Combine them to make
prediction (similar to KNN,
majority vote or average)


R Teamwork (Bagging)

Bootstrap + mtry hyperparameter


All features




Tree 1



Tree 2



Tree 3



Tree 100





Random Forest Code

```
## train random forests
set.seed(99)
myGrid <- expand.grid(mtry = 2:4)

ctrl <- trainControl(method = "...", number = ...,
                     verboseIter = TRUE)

rf_model <- train(diabetes ~ .,
                  data = ...,
                  method = "rf",
                  tuneGrid = myGrid,
                  trControl = ...)

## test model
p <- predict(rf_model, newdata = ...)
accuracy <- mean(...)
```





Ensemble Models

Ensemble

American pronunciation ▾

aan · **saam** · bl 🔊

☐ Slow

Feedback

อาน ซาม เบิ้ล



นำโมเดลหลายๆตัวมาช่วยกัน ทำนายผล (Majority Vote)

KNN	Logistic Regression	Ridge Regression	Decision Tree	Random Forest
1	0	0	1	1



Save our models for later use

```
saveRDS(model, "model.rds")  
  
model <- readRDS("model.rds")
```

- Machine Learning is **art + science**
- Try different models (No Free Lunch)
- Choose the simpler model
- Use CV + Grid Search to fine-tune model
- Start with Regression or decision tree because **they are very fast to train**

Bootcamp Live 07 Introduction to Machine Learning

Website: <https://datarockie.com>

R

