
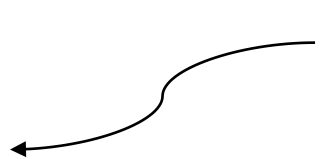

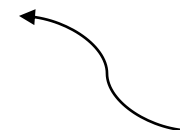


# *Introduction to* **MACHINE LEARNING**

BY DATAROCKIE 

- What is **Data Science**?  15 mins
- What is **Machine Learning**?  30 mins
- **4 Steps** in Machine Learning  30 mins
- **Algorithms** You Need to Know  60 mins

**What is  
DATA SCIENCE?**

# Data + Science

## What is Data?

Data is new electricity.

## Science as we know it.

Scientists use data to test hypothesis, theory formation.



*Andrew Ng*

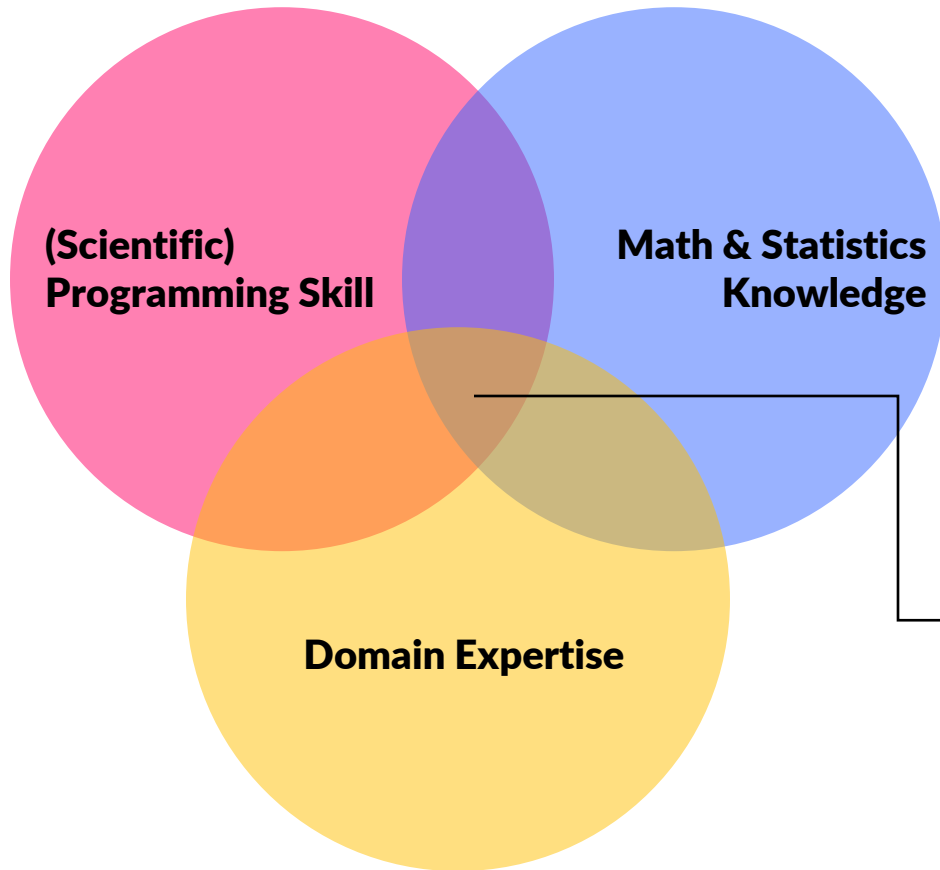
A black and white photograph of a man with a beard and short hair, wearing a dark suit jacket over a light-colored shirt. He is gesturing with his hands as if speaking. The background is a plain, light-colored wall.

“

Data is the past, and the past is Data

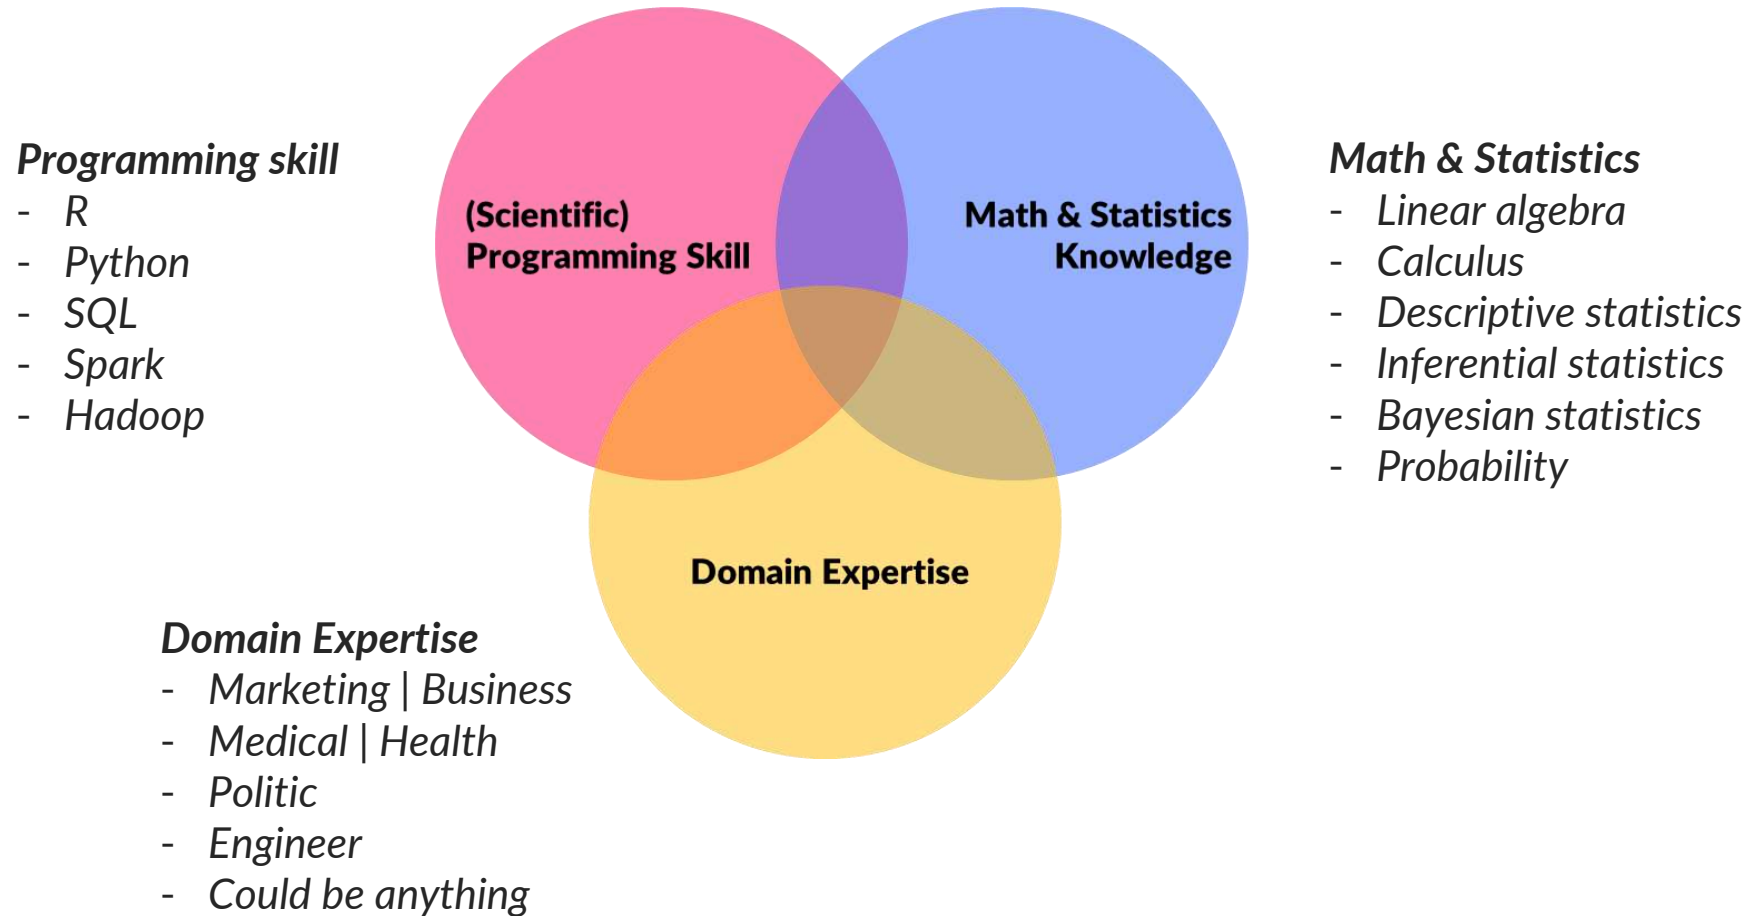
*Kirill Eremenko*

# DS Venn Diagram



**Data Science**  
*is the intersection of three skills*

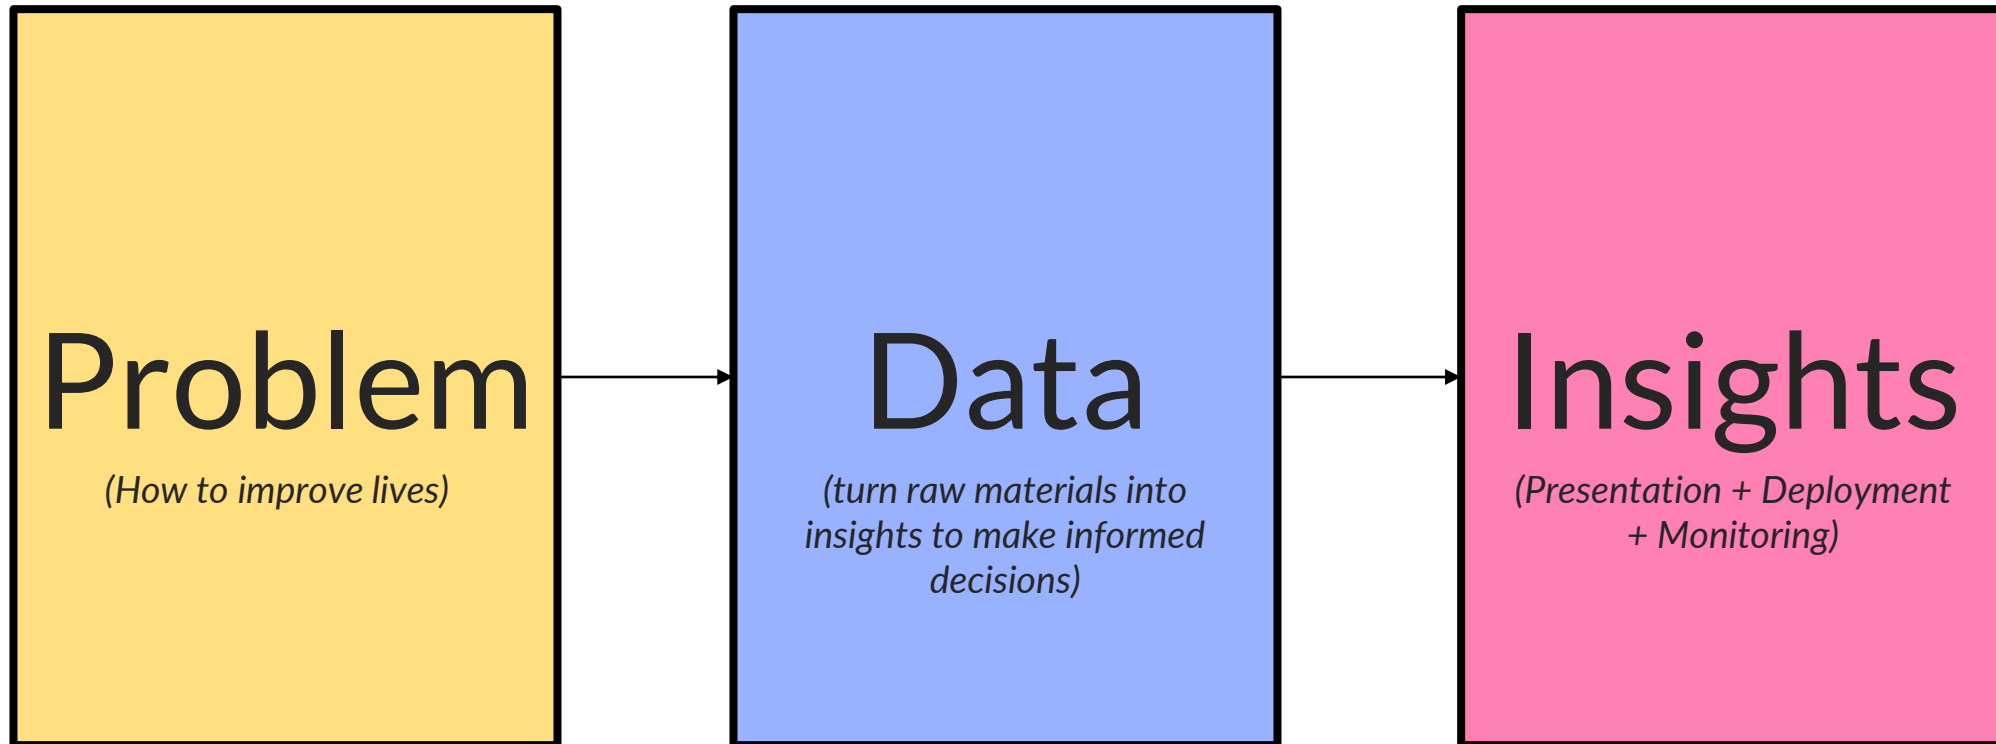
# Essential Knowledge





Always

# Good data science starts with good question





**What is  
MACHINE LEARNING?**

# Can you guess?

x (input)	1	2	3	4	5	6	7	8	9	10
y (output)	1	4	9	16	25	36	49	64	81	<div></div>

You guess the function is  $x^2$





Ask her for a date?



# Right! Can you guess this one?

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub   Bar	Many	Saturday	OK	Yes
Occasion 3	Pub   Bar	Few	Monday	OK	No
Occasion 4	Restaurant	Few	Friday	Bad	Yes
Occasion 5	Restaurant	Many	Friday	Good	No
Occasion 6	Pub   Bar	Few	Saturday	OK	Yes
Occasion 7	Pub   Bar	Many	Monday	Good	No
Occasion 8	Restaurant	Few	Saturday	OK	No
Occasion 9	Pub   Bar	Many	Friday	Bad	No
Occasion 10	Restaurant	Few	Friday	Good	

Yes or No?



**Machine Learning**  
is *function approximation*

# Machine Learning

is *inductive reasoning*

Learning from  
experience



# ML algorithms trying to **MAP** inputs (x's) to output (y)

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub   Bar	Many	Saturday	OK	Yes
Occasion 3	Pub   Bar	Few	Monday	OK	
Occasion 4	Restaurant	Few	Friday	Bad	
Occasion 5	Restaurant	Many	Friday	Good	
Occasion 6	Pub   Bar	Few	Saturday	OK	
Occasion 7	Pub   Bar	Many	Monday	Good	
Occasion 8	Restaurant	Few	Saturday	OK	
Occasion 9	Pub   Bar	Many	Friday	Bad	
Occasion 10	Restaurant	Few	Friday	Good	





**It's good to have lots data but ...**

**REPRESENTATION**

**is the quality we desire**

# ML Example

- Predict breast cancer
- Weather forecast
- Watson analytics
- Google flu trend
- Obama winning the US election



*Predicting crime before it happens*





**Can you get rid  
of poverty?**



A close-up, slightly blurred photograph of a laptop screen displaying JavaScript code. In the foreground, on the left, is a metal mesh pen holder containing several pens and pencils. The laptop screen shows a code editor with a dark background and syntax-highlighted code. The code includes comments and logic for checking if a rating prompt is disabled, setting a timer, and displaying the prompt. The text 'For some problems we wish there is no solution' is overlaid in a large, white, bold font on a black rectangular background at the bottom of the image.

**For some problems  
we wish there is no solution**

```
buttonText: 'Okay',  
type: 'warning',  
});  
  
async ratingPromptCheck() {  
  /* Check if the rating prompt is enabled (null or not) */  
  let ratingPromptDisabled = await Cache.get('ratingPromptDisabled');  
  
  if(!ratingPromptDisabled) {  
    /* Check if the timer is set or not */  
    let ratingPromptTimer = await Cache.get('ratingPromptTimer');  
  
    if(!ratingPromptTimer) {  
      Cache.set('ratingPromptTimer', _now() / 1000);  
    } else  
    if(ratingPromptTimer <= _now() / 1000 - ratingPromptTime + 600000) {  
      this.displayRatingPrompt();  
    }  
  }  
  
  async displayRatingPrompt() {  
    Alert.alert(  
      'May we ask for a goo rating?',  
      'Leave us a 5 star rating if you enjoy the app, thank you! (tap long screen)',  
      {  
        text: 'Not now', onPress: () => {  
          /* Reset the timer */  
          Cache.set('ratingPromptTimer', _now() / 1000);  
        }  
      }  
    );  
  }  
}
```

# A proper definition of ML

“ Machine Learning: Field of study that gives computers *the ability to learn* without being explicitly programmed.




Arthur Samuel (1959)

# A more modern def in 1998

Well posed Learning Problem: A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$ , if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ .



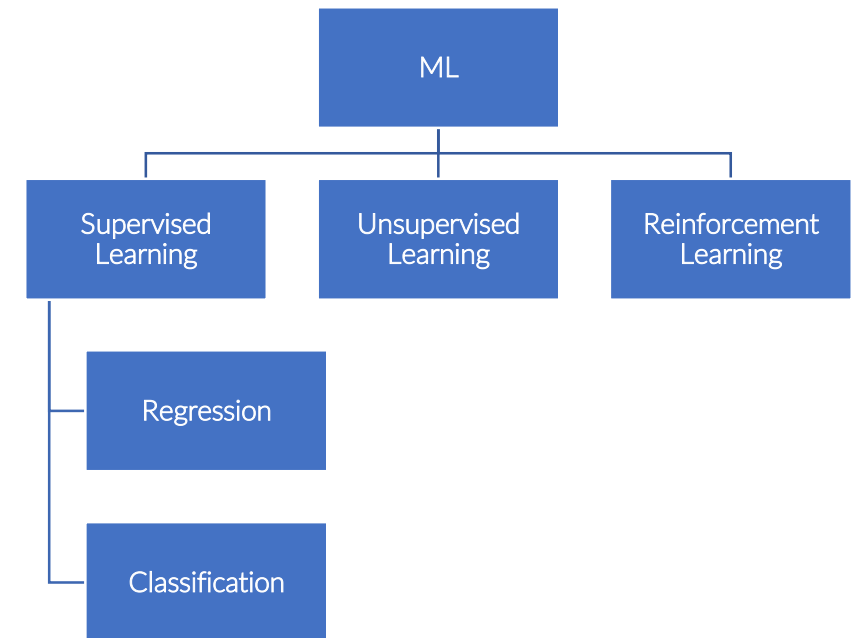
Tom Mitchell (1998)



Can we improve computer's ability to perform task  $T$  over time (without being explicitly programmed)?

# Types of ML algorithms

1. Supervised Learning
  - i. Regression
  - ii. Classification
2. Unsupervised Learning
3. Reinforcement Learning






# What's **Artificial Intelligence**



Function approximation


$$Y = f(X)$$

**Model = Algorithm (Data)**  
Predicting new unseen data In the past

*We build models to represent the world (problem we're trying to solve)*




**And the goal of Machine Learning is**

**GENERALIZATION**

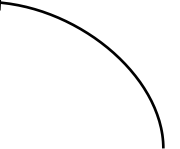
Beyond (past) data you used to build the model

**Generalization means we make sure  
our models **don't overfit** the data**


Also don't underfit




We mean  
data in the  
past




If your model **can't generalize to new unseen data**, then you're in trouble.



*Underfitting*



*Good fit*



*Overfitting*

[https://cdn-images-1.medium.com/max/1600/1\\*6vPGzBNppqMHllg1o\\_se8Q.png](https://cdn-images-1.medium.com/max/1600/1*6vPGzBNppqMHllg1o_se8Q.png)

This model looks okay as it fits underlying patterns in the data well. Though, not 100% correct?

# SUMMARY

- Data Science in Thailand is still very young
- ML is data scientist's arsenal tool
- ML is **Function Approximation**, learning from experience
- **Representation & Generalization** (avoid overfitting | underfitting)
- 3 types of ML algorithms including supervised, unsupervised and reinforcement learning
- Regression predicts numbers | Classification predicts categories

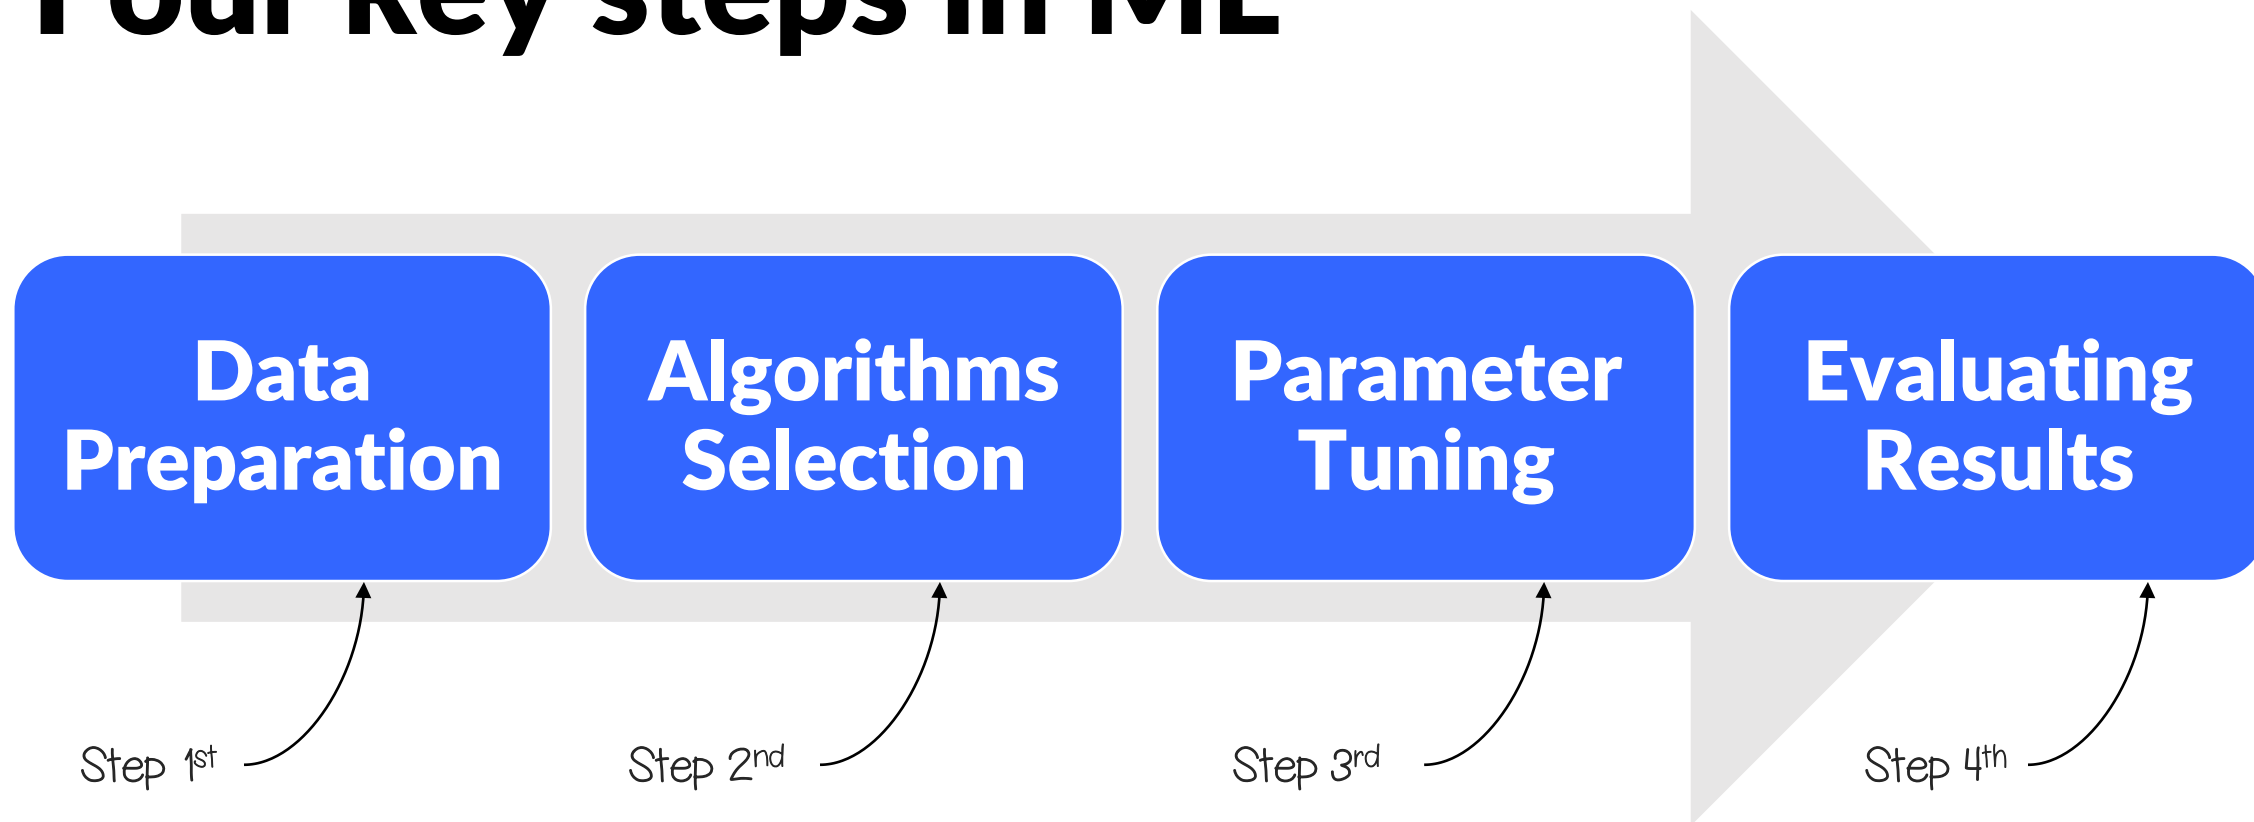


# KICK START

# MACHINE LEARNING




# Four key steps in ML



**DATA**


**PREPARATION**

# Data Format



A diagram above the table shows a long pink double-headed arrow spanning the first five columns (X1 to X4) labeled "Variables | Features" in pink text. A shorter pink double-headed arrow spans the sixth column (Y) labeled "Target Variable" in pink text.

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub   Bar	Many	Saturday	OK	Yes
Occasion 3	Pub   Bar	Few	Monday	OK	No
Occasion 4	Restaurant	Few	Friday	Bad	Yes
Occasion 5	Restaurant	Many	Friday	Good	No
Occasion 6	Pub   Bar	Few	Saturday	OK	Yes
Occasion 7	Pub   Bar	Many	Monday	Good	No
Occasion 8	Restaurant	Few	Saturday	OK	No
Occasion 9	Pub   Bar	Many	Friday	Bad	No
Occasion 10	Restaurant	Few	Friday	Good	No



A vertical pink double-headed arrow on the right side of the table, spanning all 10 rows, is labeled "Data Points | Observations" in pink text.

# Variable Types

1. Binary (1/0)
2. Categorical
3. Integer
4. Continuous

# Feature Selection

- Choose relevant set of features that represent the problem
- Curse of dimensionality:  $2^n$

# Feature Engineering

- Recode
- Combine multiple variables (PCA)
- Extract new information from existing feature

# Gentle intro to Feature Engineering

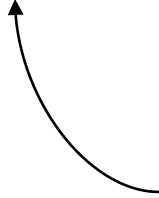
ID	Name	DOB	Math	Stats	Art	Music	Design
1	Mr. David Beckham	13 / 05 / 1975	85	90	50	42	49
2	Ms. Angelina Jolie	20 / 09 / 1988	50	60	90	95	88
3	Mr. Khalid Khan	05 / 01 / 1950	32	65	70	72	65
4	Mr. Hibino Takamoto	31 / 12 / 1995	75	76	80	50	30

Question:

**What new variable can we create from this dataset?**

**Think at least 5 new variables**

We can use  
`dplyr::mutate( )` to  
create new variables



# Missing Data

1. Approximated
2. Predicted (Supervised Learning)
3. Removed

Mean / Median  
imputation



If we have large data points and NA is not  
much, we may consider removing them





# Approximated or Removed?

```
Console ~/Desktop/R PROGRAMMING/
> head(complete.cases(biopsy), 200)
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[14] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE
[27] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[40] TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[53] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[66] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[79] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[92] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[105] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[118] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[131] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE
[144] TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[157] TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE
[170] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[183] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[196] TRUE TRUE TRUE TRUE TRUE TRUE
```

## R functions:

`complete.cases(df)` สำหรับตรวจสอบว่าข้อมูลมี NA หรือเปล่า?



`tidyr::drop_na(df)` สำหรับ drop case ที่มีข้อมูลไม่ครบทั้ง

# ALGORITHM SELECTION

# There are 3 types of algorithms

	Algorithms	Main Tasks
<b>Unsupervised learning</b>	<i>K-Means Principal Component Analysis (PCA) Association Rules Social Network Analysis (SNA)</i>	<i>Tell me what patterns exist in my data - <b>descriptive</b></i>
<b>Supervised learning</b>	<i>Linear Regression Logistic Regression Decision Tree Random Forests KNN Support Vector Machine Neural Networks</i>	<i>Use the pattern in my data to make predictions - <b>predictive</b></i>
<b>Reinforcement learning</b>	<i>Multi-Armed Bandits (UBC, Thompson Sampling)</i>	<i>Use the pattern in my data to make predictions and - <b>improve</b> - these predictions as more results come in</i>

# Regression vs. Classification



Decision Boundary:  
The line that best split  
red vs. blue data points

100% in reality may not  
be feasible, some blue  
points are mistakenly  
classified as red




A background image of a cafe interior. On the left, a menu is posted on a white tiled wall, listing items like Espresso, Cortado, Macchiato, Cappuccino, Americano, Cold coffee, Cold soy, Tea (pot), Chocolate, Black coffee, Soft drink, and Soy with their respective prices. A barista in a black shirt is behind the counter, and a customer is interacting with him. In the foreground, a woman with long dark hair and a black top is standing at the counter. To her right, another woman wearing sunglasses and a light blue shirt is standing. The background shows more of the cafe, including a large clock on the wall and shelves filled with books or pastries.

# NO FREE LUNCH


- Speed
- Accuracy
- Bias & Variance Tradeoff

Our model consists of 3 types of error

$$Y = f(X) + \text{Error}$$







# Example Models

<b>High Bias</b> <i>Having lots of assumptions to do function approximation</i>	<i>Linear Regression</i> <i>Logistic Regression</i> <i>Linear Discriminant Analysis</i>	<b>Low Variance</b> <i>New data does not cause model to change much</i>
<b>Low Bias</b> <i>Having little or few assumptions about the function approximation</i>	<i>Decision Tree</i> <i>KNN</i> <i>SVM</i>	<b>High Variance</b> <i>New data causes model to change much i.e. prediction results vary significantly compared to low variance algorithms</i>



# PARAMETER TUNING

**Music Composer** is  
a very good example




**There are so many learners.**


**Even the same learner can have  
different performances.**



# We tune parameter to get the best results **(GOOD FIT)** from our learner.



*Underfitting*



*Good fit*



*Overfitting*

[https://cdn-images-1.medium.com/max/1600/1\\*6vPGzBNppqMHllg1o\\_se8Q.png](https://cdn-images-1.medium.com/max/1600/1*6vPGzBNppqMHllg1o_se8Q.png)

The parameters that can be tuned called "**HYPERPARAMETERS**"

# Some models have **hyperparameters**, some don't.

## Stochastic Gradient Boosting

```
method = 'gbm'
```

Type: Regression, Classification

Tuning parameters:

- `n.trees` (# Boosting Iterations)
- `interaction.depth` (Max Tree Depth)
- `shrinkage` (Shrinkage)
- `n.minobsinnode` (Min. Terminal Node Size)

*Hyperparameters of GBM algorithms*

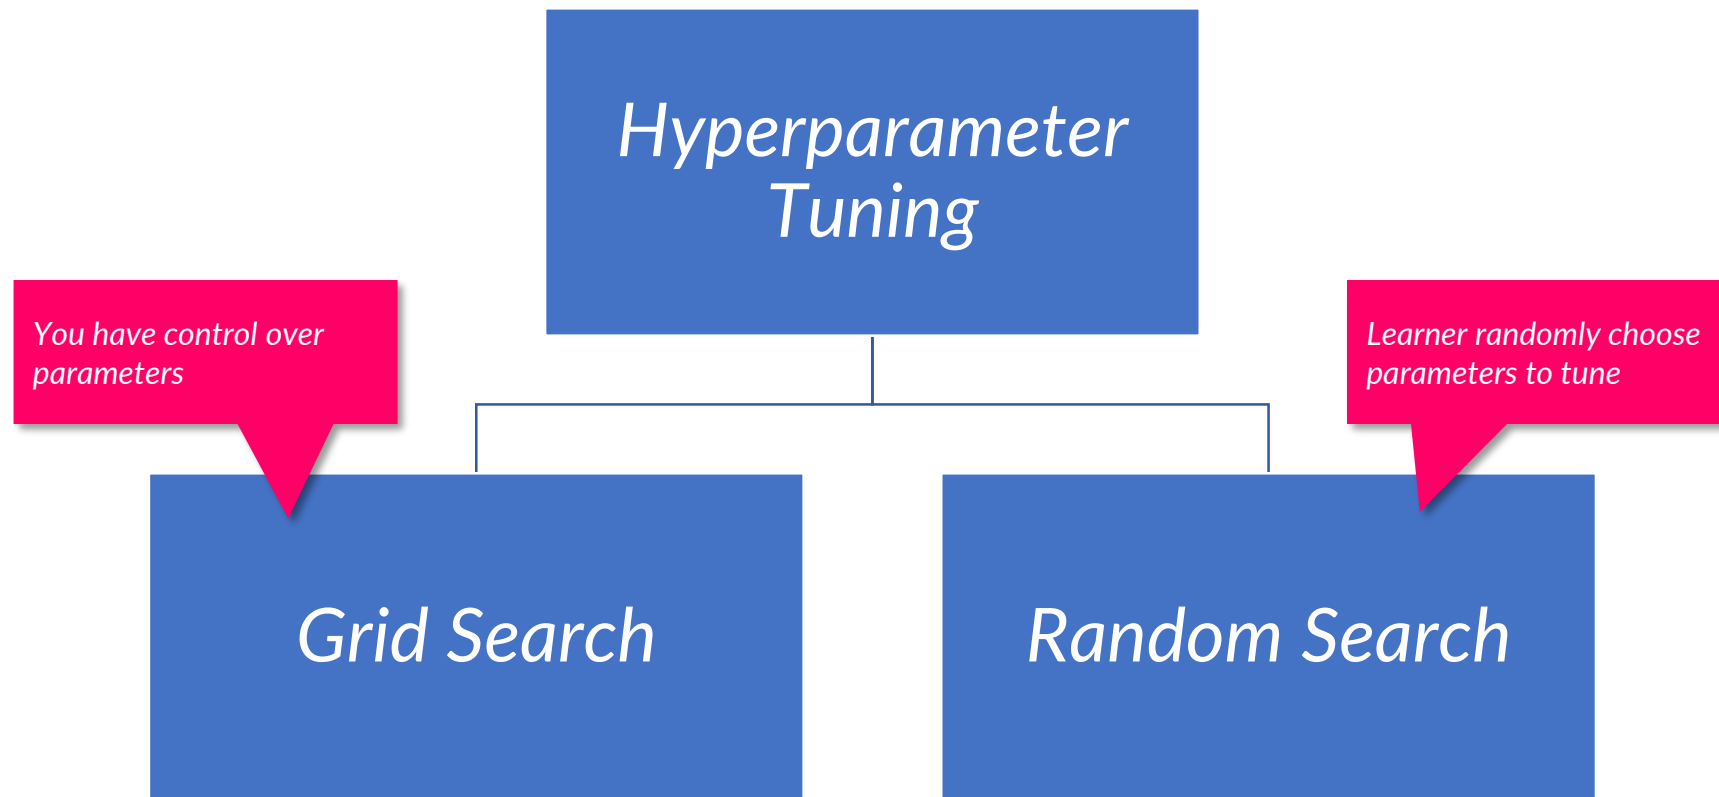
Required packages: `gbm` , `plyr`

A model-specific variable importance metric is available.

<http://topepo.github.io/caret/train-models-by-tag.html>


# SEARCH for the best tune

ML is optimization problem




It really depends on the learners. Sometimes it's more **efficient to use random search**, sometimes it's not.

Grid Layout



Important parameter

Random Layout



Important parameter



# EVALUATING RESULTS

# Performance Metrics

## Classification problems

- **Confusion Matrix**
- **% Accuracy**
- **Area Under Curve (AUC)**
- **Logarithmic Loss (Log Loss)**
- **Cohen's KAPPA** *#for imbalanced problems*

## Regression problems

- **Mean Absolute Error (MAE)**
- **Root Mean Squared Error (RMSE)**

# Cancer Detection



# Confusion Matrix

*Total n = 100*

Predicted	Reference	
	Cancer	No Cancer
Cancer	60	3
No Cancer	2	35

*Accuracy is the sum of diagonal divided by total sample size*

$$60 + 35 / 100 = 95\%$$

# Confusion Matrix

<i>Total n = 100</i>	Reference	
	Cancer	No Cancer
Predicted		
Cancer	60 (A)	3 (B)
No Cancer	2 (C)	35 (D)

ถ้าผู้ป่วยเป็นมะเร็ง เราทายถูกเท่าไร? (%)

$$\text{Sensitivity} = A / A + C$$

ถ้าผู้ป่วยไม่ได้เป็นมะเร็ง เราทายถูกเท่าไร? (%)

$$\text{Specificity} = D / B + D$$

ผล prediction "cancer" ของเราถูกกี่ %

$$\text{Precision} = A / A + B$$

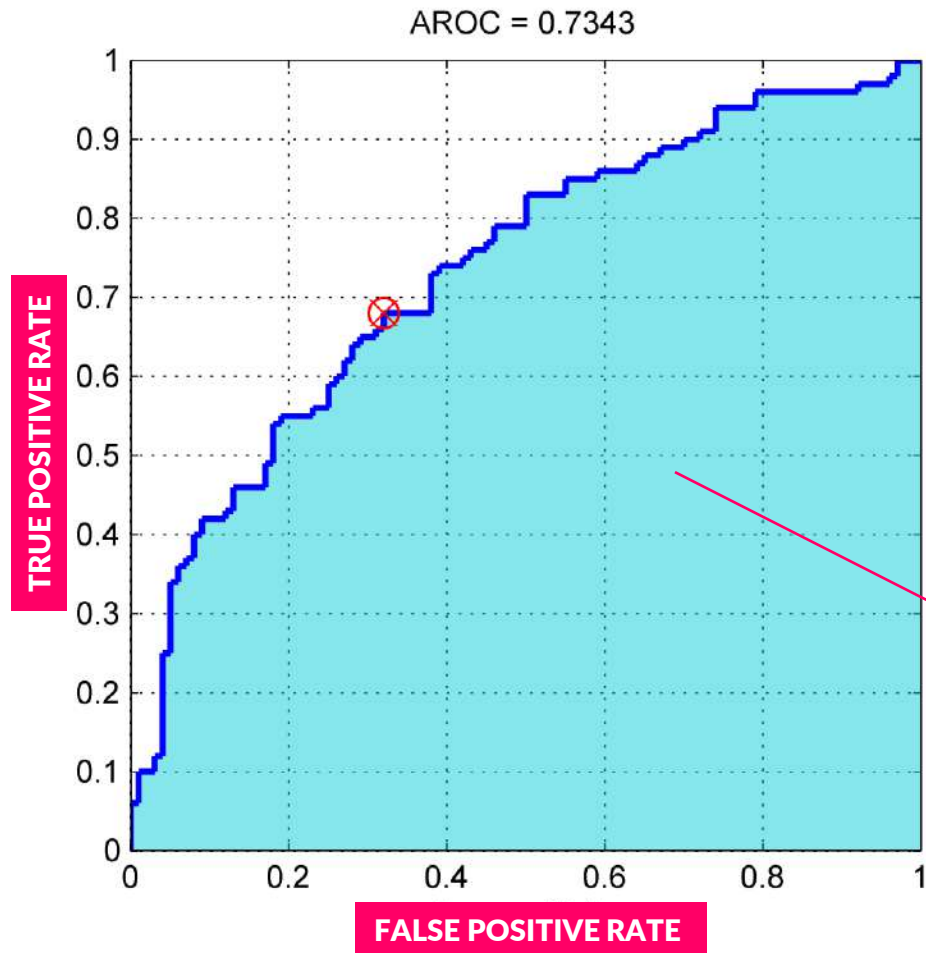
ในผู้ป่วยโรคมะเร็งทั้งหมด เราทำนายถูกเท่าไร? (%)

$$\text{Recall} = A / A + C$$

Measure of overall test's accuracy


$$F1 = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

# Area Under Curve (AUC)



*For binary classification, sometimes we want to maximize AUC*

# Root Mean Squared Error (RMSE)



Mathematical Expression always look scary

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (Actual - Predicted y_i)^2}{N}}$$

$$RMSE = \sqrt{\frac{sum(error^2)}{N}}$$

**Remember this one**


$$RMSE = \sqrt{mean(error^2)}$$



**Recall the goal of ML is**  
**“GENERALIZATION”**

**We must test our model on new**  
**unseen dataset**

# Is our model good to go?



*Training set (build model)*

*Testing set (evaluate model)*

## Decision Criterion:


*The model **passes the test** if tested performance is similar trained performance [metric]*

# Resampling techniques

1. Train-test split
2. *k-Fold Cross-Validation*
3. *Repeated k-Fold Cross-Validation*
4. *Bootstrap*
5. *Leave One Out Cross-Validation*

*This method is preferred  
but it's very expensive for  
large dataset*

# The most efficient method is train-test split



# K-Fold Cross Validation

*FULL DATASET*

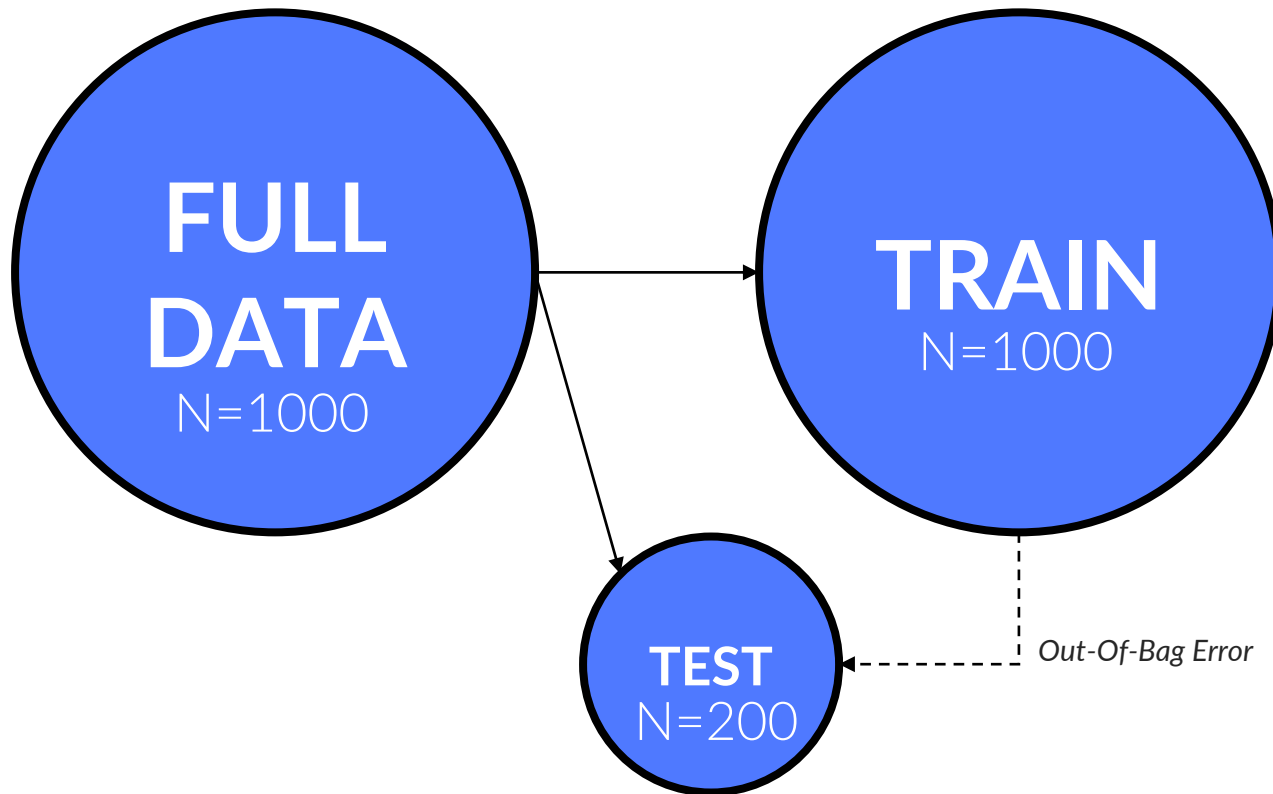


[1]	Train	Train	Train	Train	Test
[2]	Train	Train	Train	Test	Train
[3]	Train	Train	Test	Train	Train
[4]	Train	Test	Train	Train	Train
[5]	Test	Train	Train	Train	Train

People normally  
do 5, 10 K-Fold

# Bootstrap Sampling

(aka. Random sampling with replacement)

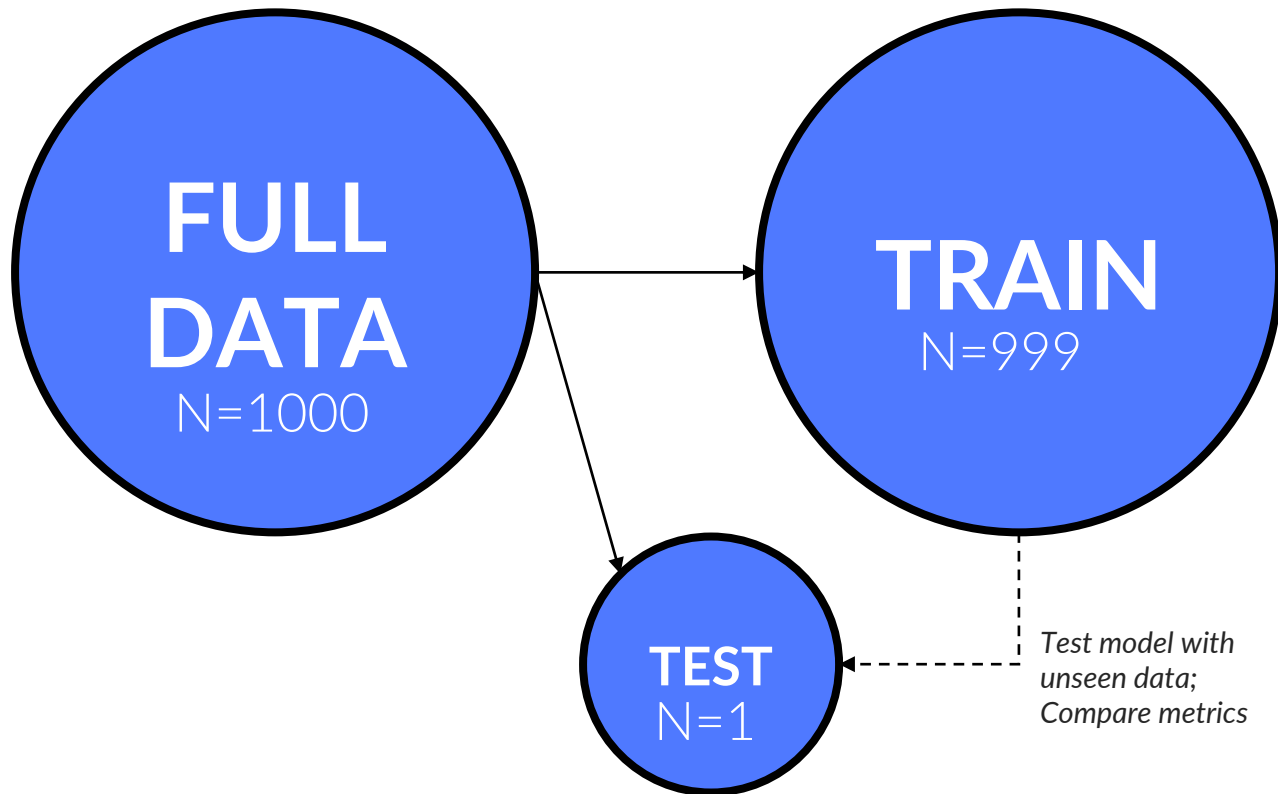


## **Explanation:**

*1000 may come from 800  
unique data points + 200  
repeated data points*

# Leave One Out CV

This is the most expensive method – time consuming

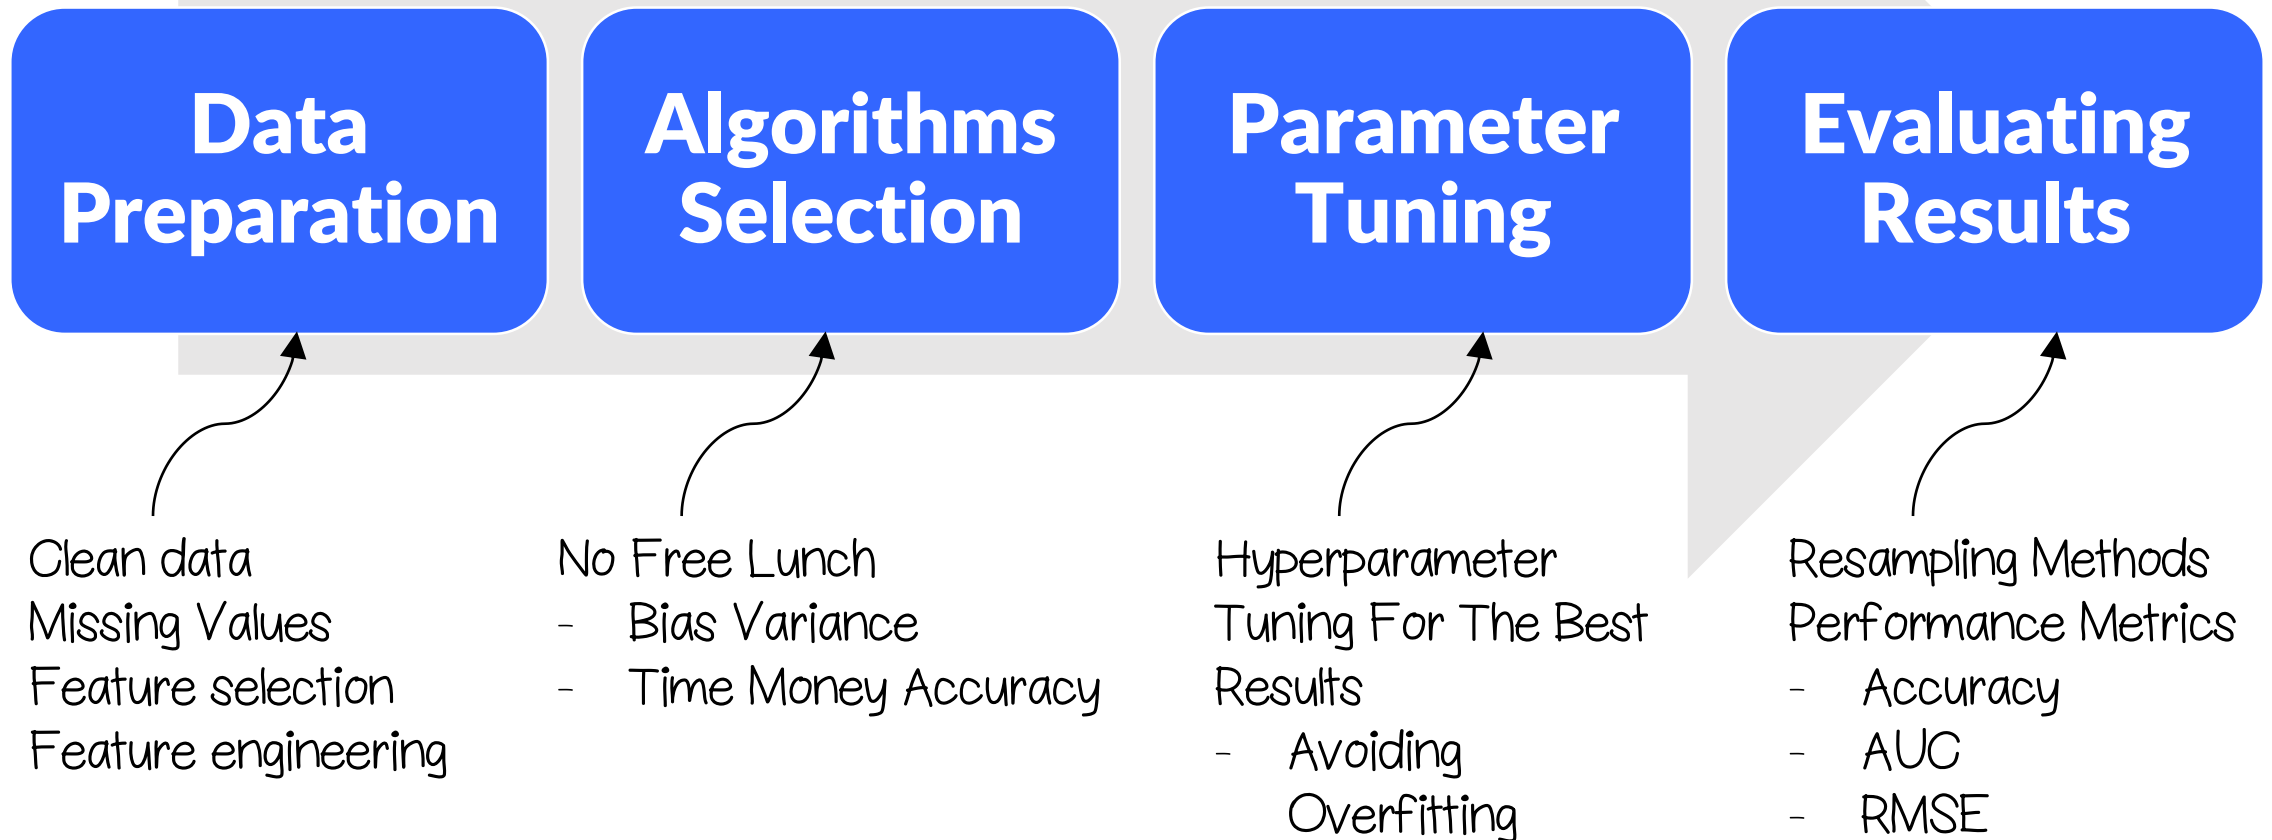


## **Explanation:**

Leave 1 data point to test the model. For  $n=1000$ , this means we need to run model 1000 times



# #SIMPLE :)



# SUMMARY

- Data preparation takes 70-80% of our time.
  - “Garbage in, Garbage out”
  - Feature selection | Feature engineering | Missing Value
- No free lunch theorem
  - **Bias vs. Variance Tradeoff**
  - **Speed vs. Accuracy Tradeoff**
- [Hyper]Parameter tuning to get the best out of our learning algorithms
- Resampling techniques used in ML research incl. **train-test, k-fold CV, repeated k-fold CV**, bootstrap, and LOOCV

# OKAY !! Let's recap a bit

- ML is function approximation
- Learn from past data to predict future (unseen) data point
- The goal is representation & generalization (avoid overfitting)
- Regression or Classification problems depend target variable ( $y$ )
- No Free Lunch: bias vs. variance tradeoff
- Algorithms can be tuned via hyperparameter
- K-Fold cross validation is popular resampling technique
- Our model passes the test if training vs. testing performances are similar (not overfit)

# ALGORITHMS


## YOU SHOULD KNOW

# We cover 8 algorithms

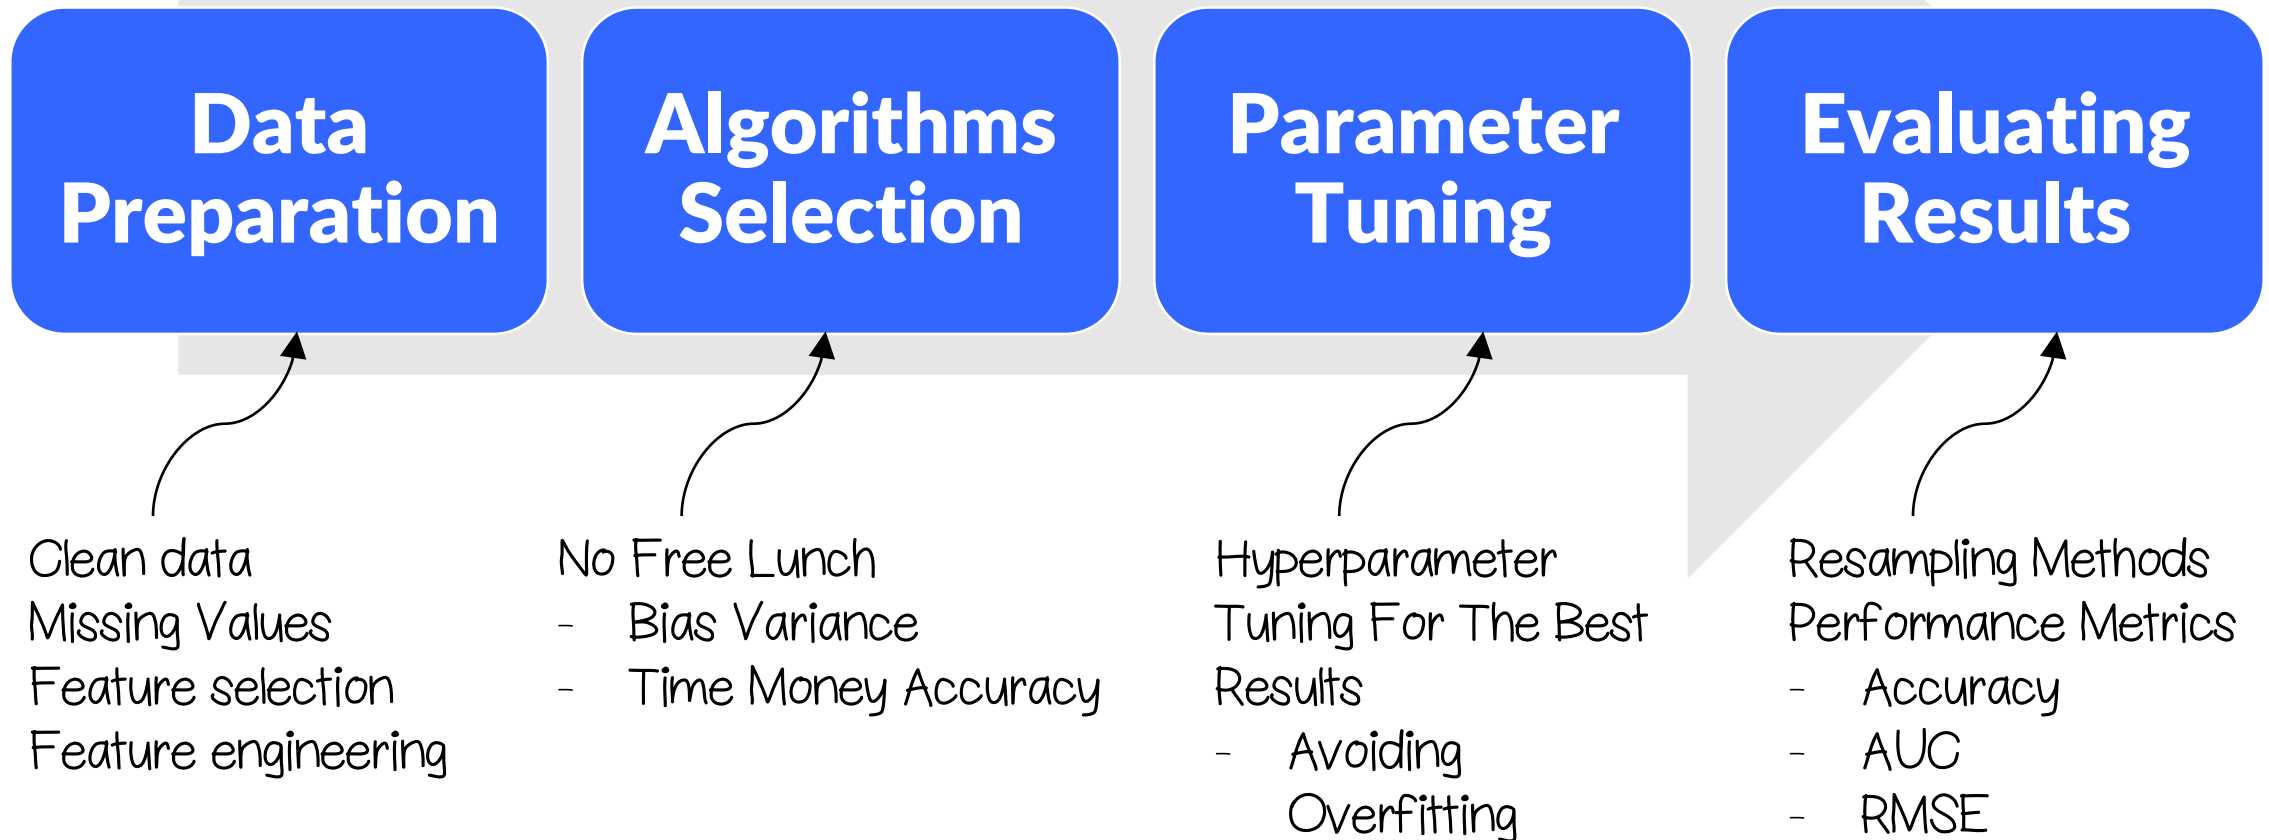
Linear Regression	Regression	Supervised Learning
Logistic Regression	Classification	
Decision Tree	Regression & Classification	
Random Forest	(mainly) Classification	
KNN	(mainly) Classification	
Support Vector Machine	(mainly) Classification	
Neural Networks	Regression & Classification	NN is special algorithm, it works like reinforcement learning
K-Means	Unsupervised Learning	

# LINEAR REGRESSION

Hello World in  
Machine Learning




# Recall these steps



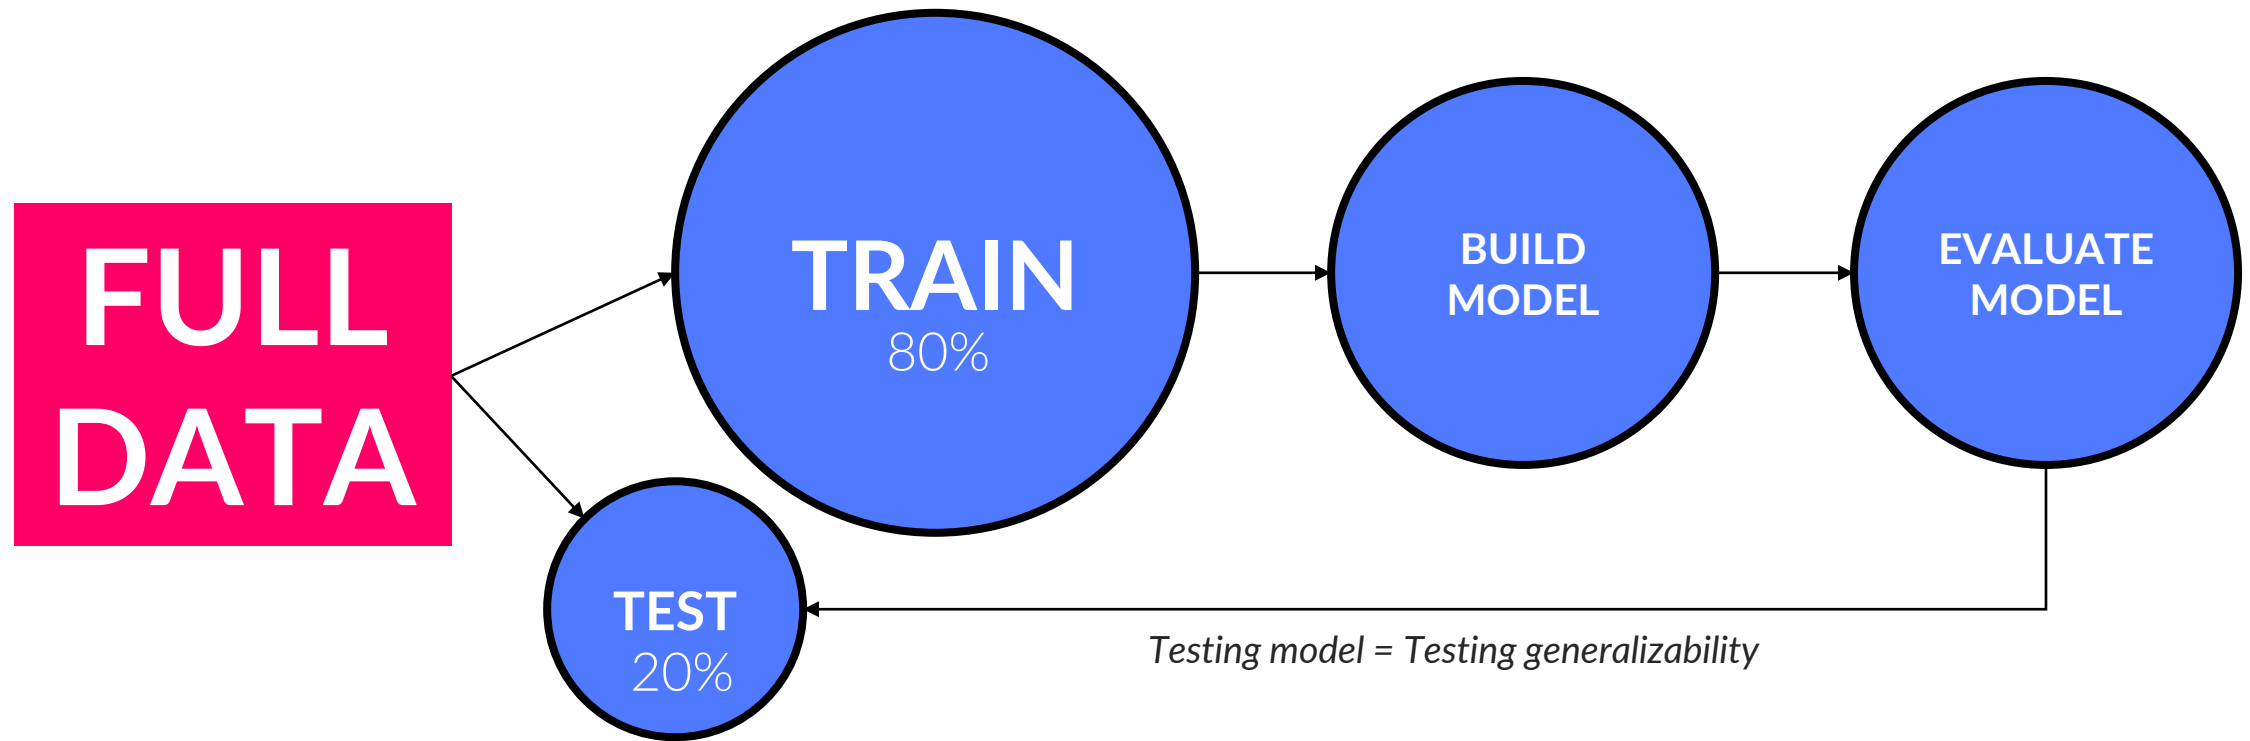


# Algorithm explained:

How to draw a single straight line that fit best with given data points



# Let's do a **train-test split** for linear regression problem



## Advantages:

- *Easy to train*
- *Always find the optimum solutions (least squared)*
- *Foundation of advanced algorithms like neural nets*


## Disadvantages:

- *Only work for regression problems (continuous  $y$ )*
- *Full of assumptions (high bias algorithms)*


# LOGISTIC REGRESSION

# Algorithm explained:

Classify data points by calculating the *probability* that  $Y=1$



# Should I get taxi or a bus?



# The model is very similar to linear regression


(assuming linear relationship between x and y)

$z = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4 \rightarrow$  then  $1 / 1 + e^{-z}$  to get probability scores

ID	X1 (minute)	X2 (crowd)	X3 (time)	X4 (weather)	P(Y=1) (get Taxi)	Prediction
1	50	Many	Friday	Good	.95	Yes
2	20	Many	Saturday	OK	.32	No
3	20	Few	Monday	OK	.31	No
4	25	Few	Friday	Bad	.85	Yes
5	30	Many	Friday	Good	.68	Yes
6	35	Few	Saturday	OK	.25	No
7	12	Many	Monday	Good	.52	Yes
8	5	Few	Saturday	OK	.12	No
9	10	Many	Friday	Bad	.20	No
10	65	Few	Friday	Good	.98	Yes



We set threshold = 0.5  
If  $p(y=1) \geq 0.5$ , then we  
predicted YES get taxi!





# Advantages:

- *One of the most powerful algorithms to date*
- *Easy to train*
- *Can be developed to create complex decision boundaries (e.g.  $x^2$   $y^2$ )*

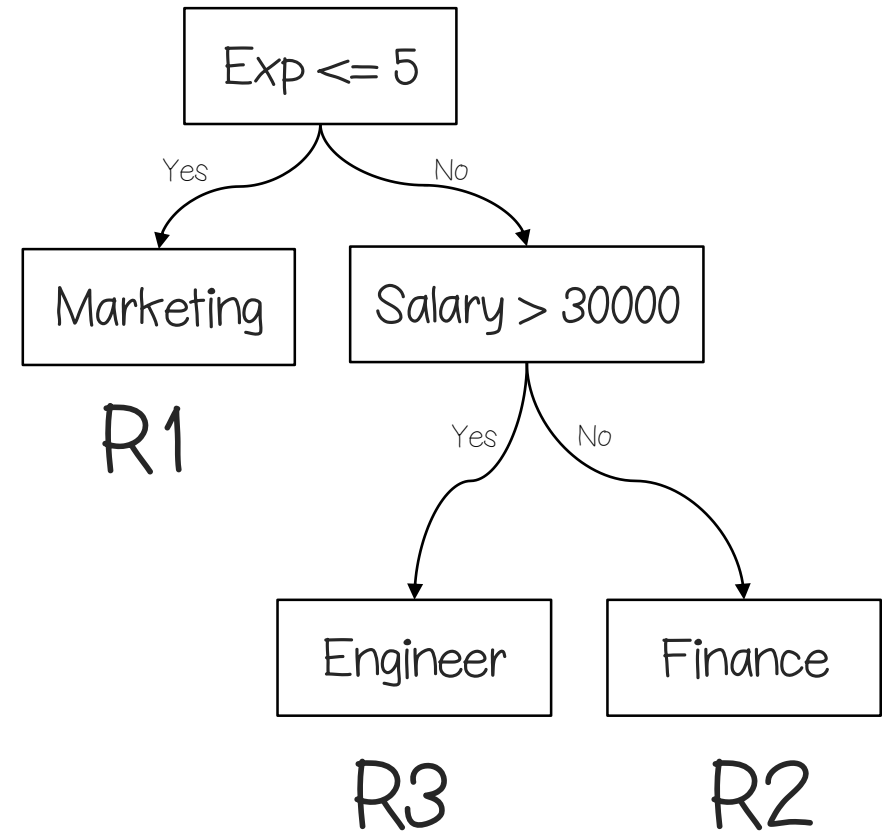

# Disadvantages:

- *Also a lot of assumption e.g. no outlier, no noise*
- *Logistic regression is also a linear algorithm*
- *Prone to multicollinearity (check correlation first)*  
*(sigmoid function transform an output into non-linear)*

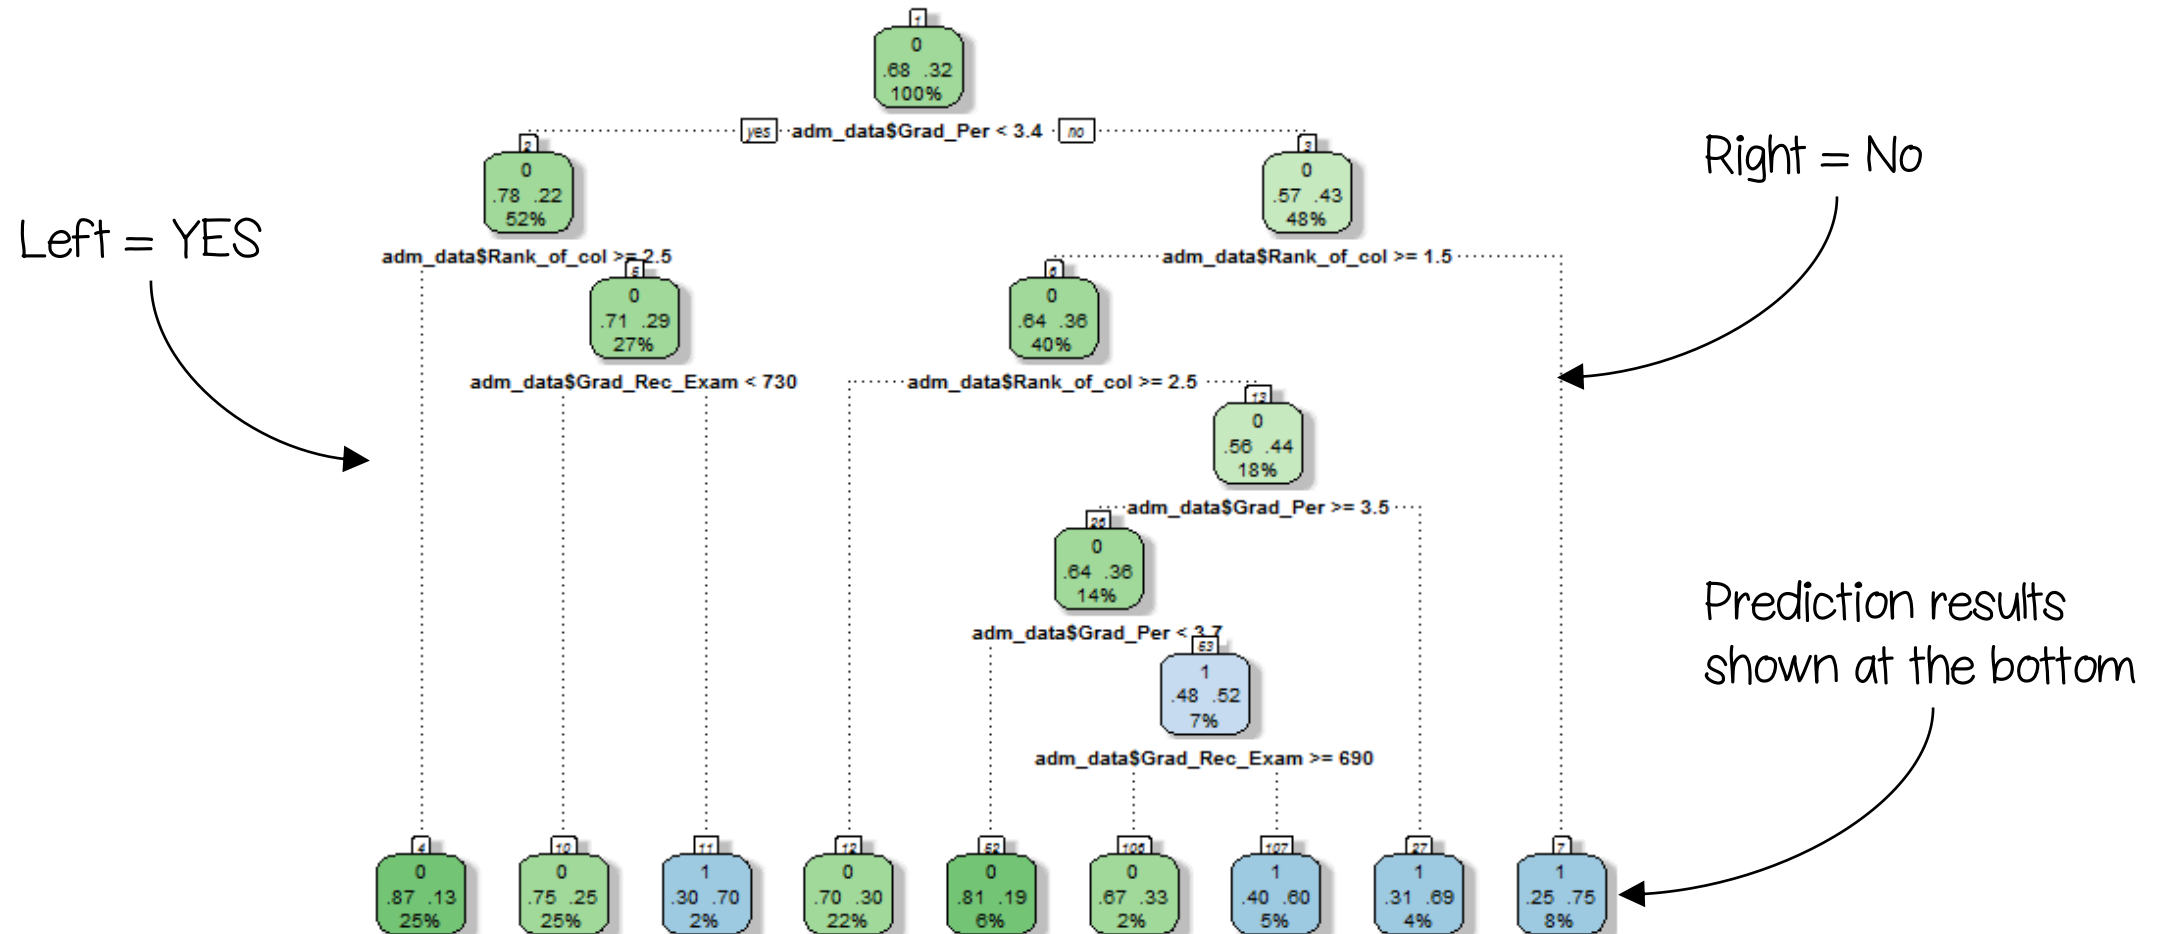
# DECISION TREE

# Algorithm explained:

Divide space into regions that best predict data points (classes)



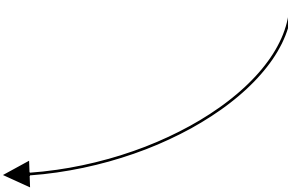
# Decision Tree is easily presented using diagram



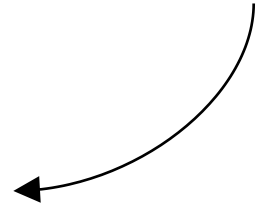
# Advantages:

- *Easy to train*
- *Easy to interpret results*
- *Widely used in many disciplines*

The main reason why people use decision tree



Locally optimal decision are made at each node, but can't guarantee globally optimal decision trees



# Disadvantages:

- *Prone to overfitting: Greedy algorithms*

[https://en.wikipedia.org/wiki/Decision\\_tree\\_learning](https://en.wikipedia.org/wiki/Decision_tree_learning)

# Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(rpart)` – train decision tree (CART): recursive partitioning and

regression tree

## CART

```
method = 'rpart'
```

Type: Regression, Classification

Tuning parameters:

- `cp` (Complexity Parameter)

Required packages: `rpart`

A model-specific variable importance metric is available.

`Cp`: complexity parameter  
to penalize larger trees




# RANDOM FOREST



# Algorithm explained:

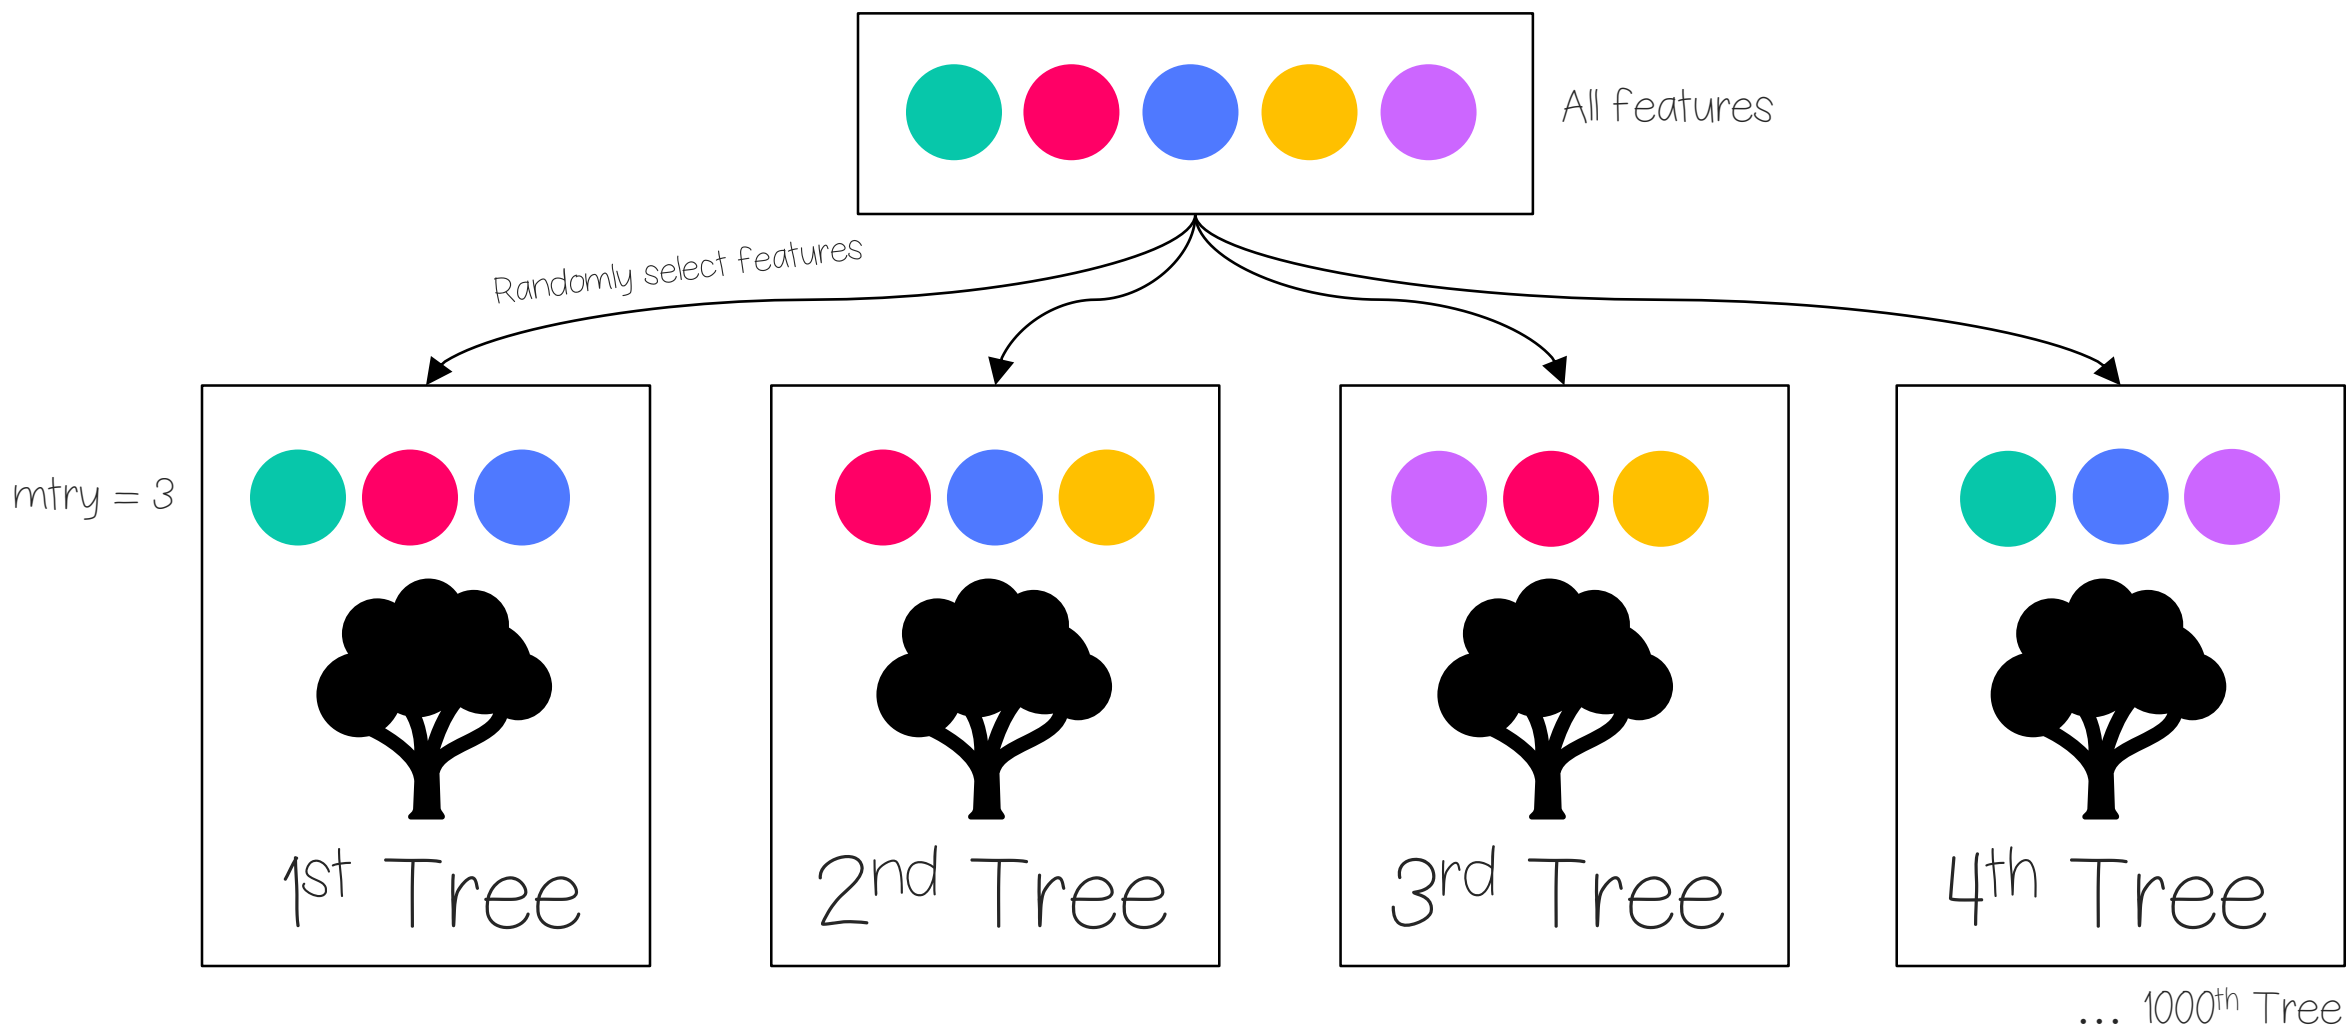
Grow thousand trees and average their prediction results



Grow *uncorrelated* 1000 trees


# How do we grow 1000 trees?

Randomness is your friend, hyperparameter is **mtry**




# Wisdom of the crowds

Find the average prediction results of all grown trees



... 1000<sup>th</sup> Tree

## Advantages:



- *Amongst the best learners in the world*
- *Easy to train (Improvement from Decision Tree)*

## Disadvantages:

- *Can be slow to train for large dataset*
- *Black Box Algorithm (not easy to explain)*

# Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(randomForest)` – train random forest in caret

## Random Forest

```
method = 'rf'
```

Type: Classification, Regression


Tuning parameters:

- `mtry` (#Randomly Selected Predictors)

Required packages: `randomForest`

A model-specific variable importance metric is available.

Tune `mtry` parameter to  
find the best prediction  
from all trees




# Let's talk about the hardware side



Graphic Cards  
have GPU 1000  
cores



CPU  
MULTIPLE CORES



GPU  
THOUSANDS OF CORES

Nvidia has CUDA technology

Read more: <https://developer.nvidia.com/about-cuda>



**KNN**

**(K-NEAREST NEIGHBOR)**




# Algorithm explained:

Predict new data according to the distance from its neighbors



# Euclidean Distance:

The most popular distance calculation method



$$d = \sqrt{\sum (x_i - y_i)^2}$$

$$d = \sqrt{(1 - 5)^2 + (6 - 2)^2}$$

$$d = \sqrt{16 + 16}$$

$$d = 5.656854$$

Very simple  
to compute

# Another Example

Consider the `nutrient` dataset provided with the `flexclust` package. The dataset contains measurements on the nutrients of 27 types of meat, fish, and fowl. The first few observations are given by

```
> data(nutrient, package="flexclust")  
> head(nutrient, 4)
```

	energy	protein	fat	calcium	iron
BEEF BRAISED	340	20	28	9	2.6
HAMBURGER	245	21	17	9	2.7
BEEF ROAST	420	15	39	7	2.0
BEEF STEAK	375	19	32	9	2.6

We can calculate many dimensions (more than two)

and the Euclidean distance between the first two (beef braised and hamburger) is

$$d = \sqrt{(340 - 245)^2 + (20 - 21)^2 + (28 - 17)^2 + (9 - 9)^2 + (2.6 - 2.7)^2} = 95.64$$

## Advantages:

- *Very simple*
- *Easy to train (remember all data)*

## Disadvantages:

- *Worst for large dataset*
- *Not good with high dimensional data (many features)*
- *Don't work well with categorical data*

# Example Code:

```
5 ## load library
6 library(caret)
7
8 ## train model
9 ## define trControl
10 my_trcontrol = trainControl(method = "cv",
11                             number = 3,
12                             verboseIter = TRUE,
13                             search = "random")
14
15 knn_model <- train(Species ~ . , data = m,
16                   method = "knn",
17                   metric = "Accuracy",
18                   tuneLength = 10,
19                   trControl = my_trcontrol)
```


Define our technique  
(to train the model)

Train the model using  
random search +  
tuneLength = 10


# SUPPORT VECTOR MACHINE

# Algorithm explained:


We want to maximize the *margin* from decision boundary



# But can you divide this with a single straight line?




Higher  
Dimension






# Impressive Result!!





Decision  
Boundary



# Let's look closely at kernel trick

- we have  $(x,y)$  coordinates
- we project a new higher dimension:  $(x, y, \underbrace{x^2 + y^2}_{3^{\text{rd}} \text{ dimension}})$



[https://en.wikipedia.org/wiki/Kernel\\_method](https://en.wikipedia.org/wiki/Kernel_method)

# Advantages:

- *Work with Regression | Classification problems*
- *Superior Performance*
- *Having lots of available kernel tricks*

<https://topepo.github.io/caret/train-models-by-tag.html#support-vector-machines>

# Disadvantages:

- *Expensive*
- *It's lazy algorithm (remember data like KNN)*


# NEURAL NETWORK



Neuron



# How your brain cells communicate





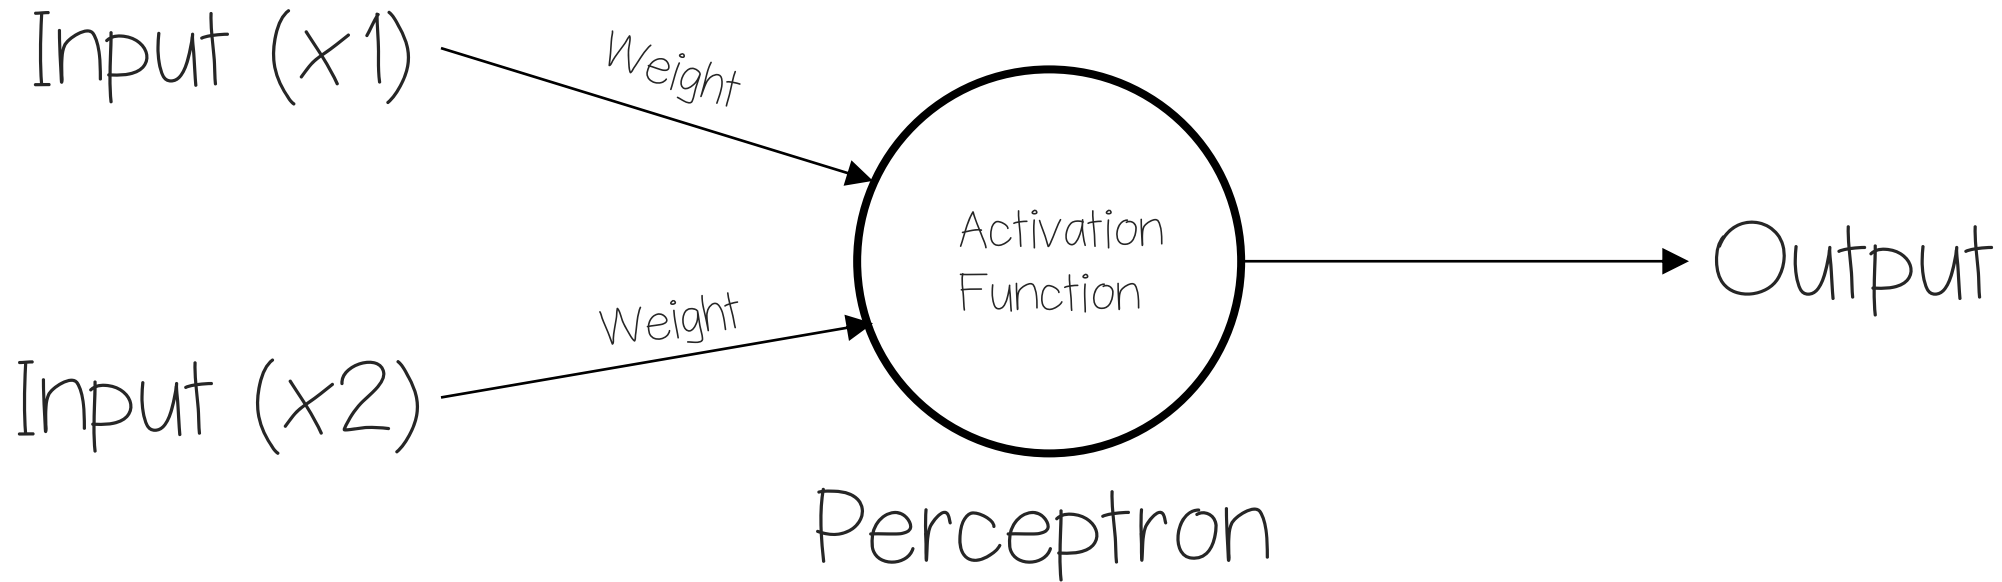


If communicated successfully,  
the neuron fire!




# Algorithm Explained:

Neurons fire only when it's the signal strong enough to activate them







Loss layer: send feedback to input layer (whether prediction is correct | wrong)


For basic neural networks, hyperparameter is number of hidden layer and learning rate (alpha)

Remarks:


NN is a good example of supervised learning that also a reinforcement learning (learn from experience)

# Deep Learning

Expands the number of hidden layers



CIFAR-10 ResNets



Hieu Pham, Has done some machine learning

Updated Apr 29, 2017




In this [paper](#) (Huang et al, 2016), the authors report that they “can increase the depth of residual networks *even beyond 1200 layers* and still yield meaningful improvements in test error (4.91% on CIFAR-10)”.

I have never heard of anything deeper than that.



# Backpropagation

The backbone of artificial neural networks



Wrong

Right

Send feedback to  
input layers and  
adjust the weights



# NO FREE LUNCH

- Speed
- Accuracy
- Interpretability
- Bias & Variance Tradeoff

ความสามารถในการ  
อธิบายผลโมเดล





# Advantages:

- Amongst the very best performers | learners in the world
- Can be developed into deep learning algorithms
- Work with both regression and classification problems

# Disadvantages:

- Requires large data set
- Black Box
- Prone to overfitting problem
- Need more power if more hidden layers

Time to buy a  
graphic card :D



# ALGORITHMS ARE NOT CREATED EQUAL

- **Neural Nets** [Backpropagation]
- **SVM** [Kernel Tricks]
- **KNN** [Majority Vote]
- **Random Forest** [Randomness]
- **Decision Tree** [Recursive Partitioning]
- **Logistic Regression** [Sigmoid Function]
- **Linear Regression** [Least Squared]



*R has powerful package written by  
Max Kuhn (10 years in development)*

**# CARET**

**# Classification And Regression Training**

**# 238 models can be trained using caret**

```
install.packages("caret")
```

```
library(caret)
```

*For more details, go to this website*

<http://topepo.github.io/caret/index.html>