

## 1) Connected Components

**Your Output for 4-connected Labeling:**

	1	1		2		3	3	3		4	4	4
5		1		2		3	3	3		4		4
5	5		6	2		3	3	3		4		4
	5		6		7	3		3	3			
	5	5		8	7	3		3	3	3		9
10		5		8	7	3						9
10	10			8	7	3	3	3	3	3		9
10	10	10		8			3					
				8	8	8		11	11	11		12
13	13	13		8		8		11		11		12
13	13	13	13	8		8	8	8		11	11	11
						8	8	8		11	11	11

Equivalence Table

2	6
3	7
7	8
8	11
8	13
11	12

**Your Output for 8-connected Labeling:**

1	1			2		3				4				5
1		1		2		3				4		6		5
1			1			3				4		6		5
	1		1		3			7			4	4	4	
	1			1				7				4		
		1		1				7	7	7	4	4		
			1	1								4		
				1		8	8		9	9	4	4	4	
	10	10	1	1				8				4		
10				1		11		8		12		4		13
10						11				12				13
10	10	10	10	10	10	10				12				13

Equivalence Table

1	2
4	6
1	3
4	7
4	9
1	10
8	9
10	11

## 2) Handwritten Digit Recognition

Link: [https://colab.research.google.com/drive/11Ni0mylkZ4DjyC1L8GOIz5tG6g7iEA\\_R?usp=sharing](https://colab.research.google.com/drive/11Ni0mylkZ4DjyC1L8GOIz5tG6g7iEA_R?usp=sharing)

Source Code

```
[207] import numpy as np
import cv2
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from skimage.color import rgb2gray
from skimage.util import img_as_float
from skimage.feature import hog
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
```

```
[208] from google.colab.patches import cv2_imshow
def my_imshow(title, img ):
    print(title)
    cv2_imshow(img)
```

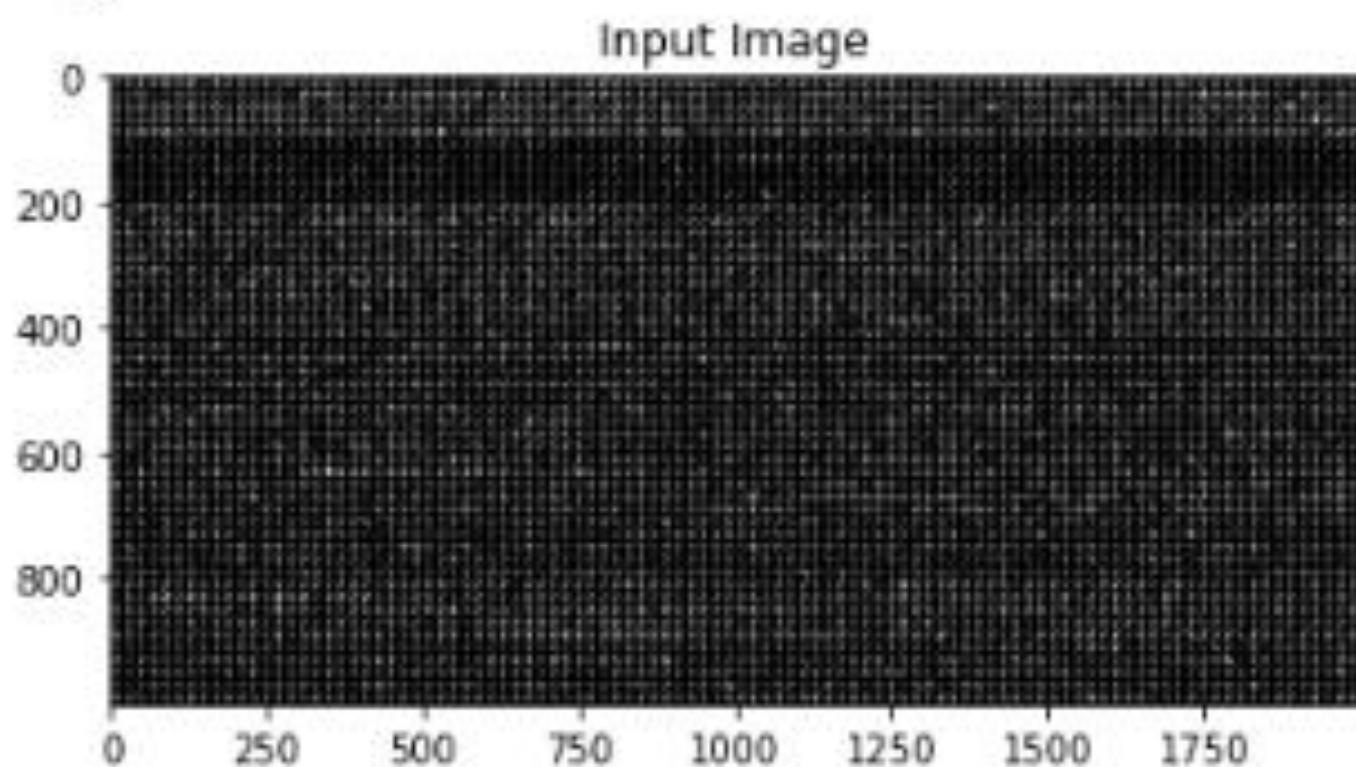
```
[224] # Load the image containing the digits
# img = cv2.imread('digits.png', 0)
path = ""
fileName = path + "digits.png"
img = cv2.imread(fileName, cv2.IMREAD_GRAYSCALE)

# Print error message if image is null
if img is None:
    print('Could not read image')
else:
    print("Image file read success...")

plt.title("Input Image")
plt.imshow(img, cmap='gray')
plt.show()
```

Output

Image file read success...



```
[215] # Reshape the image to extract each digit
digits = np.array([np.hsplit(row, 100) for row in np.vsplit(img, 50)])
digits = digits.reshape(-1, 20, 20)

# Rescale from 20 x 20 to 24 x 24
digits_rescaled = []
for digit in digits:
    rescaled_digit = cv2.resize(digit, (24, 24))
    digits_rescaled.append(rescaled_digit)


# Convert the digit images to a numpy array and normalize the pixel values
digits_rescaled = np.array(digits_rescaled, dtype=np.float32)

# Define the labels for each digit
labels = np.repeat(np.arange(10), len(digits)/10)

# Split the data into training and testing sets
training_features, testing_features, training_labels, testing_labels = train_test_split(digits_rescaled, labels, test_size=0.2, random_state=42)

# Visualization a sample digit
# Examples can be changed at training_features[i] and training_labels[i], Start i=0
for i in range(2):
    plt.title("Label:" + str(training_labels[i]))
    plt.imshow(training_features[i], cmap='gray')
    plt.show()
    cv2.imwrite('output{}.png'.format(i+1), training_features[i])
```

## Output



a)

```
[216] # Extract features from the images
def extract_gray_scale_features(X):
    return np.mean(X, axis=(1,2)).reshape(-1, 1)

def extract_HOG_features(data):
    features = []
    for image in data:
        # convert the image to grayscale
        if len(image.shape) == 3 and image.shape[-1] == 3:
            gray_image = rgb2gray(image)
        else:
            gray_image = img_as_float(image)
        # extract the HOG features
        hog_features = hog(gray_image, orientations=20, pixels_per_cell=(4, 4), cells_per_block=(2, 2), block_norm='L2-Hys', visualize=False, feature_vector=True)
        features.append(hog_features)
    return np.array(features)

X_train_gray = extract_gray_scale_features(training_features)
X_train_HOG = extract_HOG_features(training_features)
X_test_gray = extract_gray_scale_features(testing_features)
X_test_HOG = extract_HOG_features(testing_features)

# Train KNN classifiers
knn_5_gray = KNeighborsClassifier(n_neighbors=5).fit(X_train_gray, training_labels)
knn_5_HOG = KNeighborsClassifier(n_neighbors=5).fit(X_train_HOG, training_labels)
knn_1_gray = KNeighborsClassifier(n_neighbors=1).fit(X_train_gray, training_labels)
knn_1_HOG = KNeighborsClassifier(n_neighbors=1).fit(X_train_HOG, training_labels)

# Test classifiers and report accuracy
def test_classifier(classifier, X_test, y_test):
    y_pred = classifier.predict(X_test)
    accuracy = np.mean(y_pred == y_test)
    return accuracy

print("KNN K=5, gray scale features accuracy: %5.2f%%" % (test_classifier(knn_5_gray, X_test_gray, testing_labels)*100))
print("KNN K=5, HOG features accuracy: %5.2f%%" % (test_classifier(knn_5_HOG, X_test_HOG, testing_labels)*100))
print("KNN K=1, gray scale features accuracy: %5.2f%%" % (test_classifier(knn_1_gray, X_test_gray, testing_labels)*100))
print("KNN K=1, HOG features accuracy: %5.2f%%" % (test_classifier(knn_1_HOG, X_test_HOG, testing_labels)*100))
```

## Output

KNN K=5, gray scale features accuracy: 18.00%

KNN K=5, HOG features accuracy: 94.50%

KNN K=1, gray scale features accuracy: 14.70%

KNN K=1, HOG features accuracy: 96.00%

b)

ข้อ b.i

```
[217] # Create a KNN classifier with k = 1 and HOG features
      knn_1_HOG = KNeighborsClassifier(n_neighbors=1).fit(X_train_HOG, training_labels)
      print("KNN K=1, HOG features accuracy: %5.2f%%" % (test_classifier(knn_1_HOG, X_test_HOG, testing_labels)*100))


[256] # Rescale and compute HOG features for the test images
      # Path for test images
      test_image_path = "/content/drive/MyDrive/KMUTT/CPE393_ML/Quiz3/"

      # Rescale and compute HOG features for the test images
      test_image_files = ["0d1.jpg", "0d2.jpg", "1d1.jpg", "1d2.jpg", "2d1.jpg", "2d2.jpg", "3d1.jpg", "3d2.jpg", "4d1.jpg", "4d2.jpg",
                          "5d1.jpg", "5d2.jpg", "6d1.jpg", "6d2.jpg", "7d1.jpg", "7d2.jpg", "8d1.jpg", "8d2.jpg", "9d1.jpg", "9d2.jpg"]
      test_images = []
      for filename in test_image_files:
          image = cv2.imread(test_image_path + filename, cv2.IMREAD_GRAYSCALE)
          if image is not None:
              resized_image = cv2.resize(image, (24, 24))
              hog_features = hog(resized_image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
              test_images.append((resized_image, hog_features))
              plt.title("File name: " + filename)
              plt.imshow(resized_image, cmap='gray')
              plt.show()
          else:
              print(f"Could not read file {filename}")
```

## Output

KNN K=1, HOG features accuracy: 96.00%

แสดงเป็นเลข 0-9





ข้อ b.ii

```
[254] # Rescale and compute HOG features for the test images
# Path for test images
test_image_path = "/content/drive/MyDrive/KMUTT/CPE393_ML/Quiz3/"

# Rescale and compute HOG features for the test images
test_image_files = ["0p1.jpg", "0p2.jpg", "1p1.jpg", "1p2.jpg", "2p1.jpg", "2p2.jpg", "3p1.jpg", "3p2.jpg", "4p1.jpg", "4p2.jpg",
                    "5p1.jpg", "5p2.jpg", "6p1.jpg", "6p2.jpg", "7p1.jpg", "7p2.jpg", "8p1.jpg", "8p2.jpg", "9p1.jpg", "9p2.jpg"]
test_images = []
for filename in test_image_files:
    image = cv2.imread(test_image_path + filename, cv2.IMREAD_GRAYSCALE)
    if image is not None:
        resized_image = cv2.resize(image, (24, 24))
        hog_features = hog(resized_image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
        test_images.append((resized_image, hog_features))
        plt.title("File name: " + filename)
        plt.imshow(resized_image, cmap='gray')
        plt.show()
    else:
        print(f"Could not read file {filename}")
```

## Output

แสดงเป็นเลข 0-9




ข้อ b.iii

```
[12] true_labels = [5,5,5,5]
test_digits = []
test_hog_features = []
test_labels = []

# Rescale and compute HOG features for the test images
# Path for test images
test_image_path = "/content/drive/MyDrive/KMUTT/CPE393_ML/Quiz3/"

# Rescale and compute HOG features for the test images
test_image_files = ["5b1.jpg", "5b2.jpg", "5b3.jpg", "5b4.jpg"]
for filename, true_label in zip(test_image_files, true_labels):
    image = cv2.imread(test_image_path + filename, cv2.IMREAD_GRAYSCALE)
    if image is not None:
        image_resized = cv2.resize(image, (24, 24))
        test_digits.append(image_resized)
        hog_feature = hog(image_resized, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
        test_hog_features.append(hog_feature)
        test_labels.append(true_label)
        plt.title("File name: " + filename)
        plt.imshow(image_resized, cmap='gray')
        plt.show()
    else:
        print(f"Could not read file {filename}")
test_hog_features = np.array(test_hog_features)
test_labels = np.array(test_labels)
```

## Output





c)

```
[20] # Set the path to the directory containing the digits.png file
path = ""

# Load the dataset
fileName = path + "digits.png"
img = cv2.imread(fileName, cv2.IMREAD_GRAYSCALE)
if img is None:
    print("Could not read image")
else:
    print("Image file read success...")
    plt.title("Input Image")
    plt.imshow(img, cmap='gray')
    plt.show()
    # Resize the image to a uniform size of 5000 pixels per digit
    digit_size = 5000
    digits = cv2.resize(img, ((digit_size + 1) * 10, (digit_size + 1) * 5))

# Split the resized image into individual digit images and store them in an array
digits = [digits[y:y+digit_size, x:x+digit_size] for y in range(0, (digit_size + 1) * 5, digit_size + 1) for x in range(0, (digit_size + 1) * 10, digit_size + 1)]
digits = np.array(digits)

# Assign labels to each digit image (0 to 9)
labels = np.repeat(np.arange(10), len(digits) // 10)

# Apply automatic thresholding to each digit image and find the bounding box of the digit
digit_images = []
for digit in digits:
    _, thresh = cv2.threshold(digit, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)
    contours, _ = cv2.findContours(thresh, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
    x, y, w, h = cv2.boundingRect(contours[0])
    digit_image = digit[y:y+h, x:x+w]
    digit_image = cv2.resize(digit_image, (24, 24))
    digit_images.append(digit_image)

# Compute the HOG features for each digit image
hog_features = []
for digit_image in digit_images:
    hog_feature = hog(digit_image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2))
    hog_features.append(hog_feature)
hog_features = np.array(hog_features)
labels = np.array(labels)

# Split the dataset into training and testing sets
train_hog_features, test_hog_features, train_labels, test_labels = train_test_split(hog_features, labels, test_size=0.2)


# Train a KNN classifier with K = 1 on the HOG features of the training set
knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(train_hog_features, train_labels)

# Predict the labels for the testing set using the trained KNN classifier
test_pred_labels = knn.predict(test_hog_features)

# Compute the accuracy score for the predicted labels
accuracy = 100*accuracy_score(test_labels, test_pred_labels)
print("Hog Features:", hog_features.shape)
print("Labels:", labels.shape)
print("Accuracy: %5.2f%%" % accuracy)
```

## Output

Image file read success...



Hog Features: (50, 144)

Labels: (50,)

Accuracy: 10.00%